

# Compilers I

Introduction to Compiler Optimization

Dr. William L. Harrison

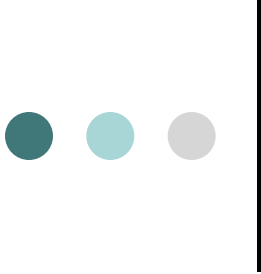
[harrisonwl@missouri.edu](mailto:harrisonwl@missouri.edu)



# Today's Class

Today: Optimizations

- Basic Optimizations
- Static Single-Assignment



# What does an optimizing compiler do?



Optimizing compiler attempts to:

- ♦ Eliminate language abstractions
- ♦ Map source program to target machine efficiently.
  - ♦ Use the hardware well!
- ♦ Equal the efficiency of a good assembly programmer.



# Dimensions of Optimization

You can optimize for a number of things.

- machine cycles (speed)
- binary size
- heap usage
- to use special hardware features
  - parallel computing

**Minimizing runtime cycles** is the typical measure to optimize.



# Proebsting's Law

- Moore's Law: Computing power doubles every 18 months.
  - This has been true for 40 years!
- Proebsting's Law: **Compiler Advances** double computing power every 18 *Years*
  - Better optimizations
  - Better use of new hardware features



# Optimizations provide engineering alternatives, but...

Imagine writing a highly optimized ray-tracer

- Lots of low-level hacks
- Difficult to maintain, extend, port, etc.

Now imagine writing a high-level mathematical specification of the ray tracer.

- Easy to modify, extend, etc.
- Let an optimizing compiler produce efficient assembly language...



YMMV



# Optimizing compilers allow efficient abstractions

- We can engineer our software applications using abstractions like OO-style objects, etc.
  - Ease of maintenance.
  - Almost all design patterns are abstractions.
- The compiler can remove much of the overheads associated with using these abstractions.



# Appel's Compiler Theorem

**Theorem:** For any optimizing compiler, there exists a better one.

*“full employment theorem for compiler writers...”*

**Appel's point was you can always add a new optimizations to make any compiler better.**



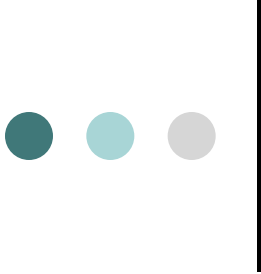
# What is an optimization?

- The output program must have the same **observable behavior**, as defined in the semantics of the source language.

```
int foo(int i) {  
    int j = 9;  
    return i + j;  
}
```

≈

```
int foo(int i) {  
    return i+9;  
}
```



# Optimizations can make your working program fail!

- The output program must have the same **observable behavior**, as defined in the semantics of the source language.
- If a program steps outside the semantically defined part of the language, all bets are off.
  - Example: using an un-initialized variable in C.  
The value might be different between optimized and non-optimized.
- Your program was faulty to start with!
  - It just happened to work...



# Optimizations

- Optimizations change the program internally
- This enables more optimizations...
  - Some optimizations are ***enabling*** optimizations
  - Some optimizations are ***true*** optimizations



# Optimization Camps

Optimizations that move computation from runtime to compile time.

- *Examples: #ifdef, asserts, constant folding*

Optimizations that replace some instructions with less expensive instructions.

- *Examples: peephole optimizations*

Optimizations that move a computation to a less expensive place.

- *Examples: loop invariant hoisting*


# Compile Time Evaluation

Sometimes  
constants can be  
evaluated at  
compile time.

## Caveats

**x** must be constant.

```
final int x = 9;  
if (x < 10) {  
    ...  
}
```



```
final int x = 9;  
if (true) {  
    ...  
}
```




# Constant folding

- Constant folding is a true optimization
  - The work that would have been done at runtime is no longer needed.
- Constant folding is **also** an enabling optimization.
  - On previous page, we can eliminate the if/then test.
  - Constant folding can give rise to **dead-code**.

# Peephole Optimizations

Consider these instructions

ADDI	$r_1 \leftarrow V_a$
ADD	$r_2 \leftarrow \text{fp} + r_1$
LOAD	$r_3 \leftarrow M[r_2 + 0]$



But

- $V_a$  is a constant.
- $\text{fp}$  is a register.

→ so we could write

LOAD $r_3 \leftarrow M[\text{fp} + V_a]$
------------------------------------------

Works inside basic blocks

Need to know about  
usage of  $r_1$  and  $r_2$ .

# Procedure inlining

## Two steps to inlining

- Substitute procedure body at call site
- Watch for name clashes
  - Fix variable names if clash

```
int foo(int x) {  
    return x * x;  
}
```

```
...  
    v := foo(y);  
...
```

Saves cost of  
function call



```
...  
    v := y * y;  
...
```



# Loop invariants

Assuming  $v$  does not change inside this loop

Can “lift” the computation of  $v * 2$  out of the loop.

```
v := ...  
...  
while (x < 10)  
{  
  y = v * 2;  
  x = x + y;  
}
```



It is almost  
always good  
to move things  
out of loops.

```
v := ...  
...  
y := v * 2;  
while (x < 10) {  
  x = x + y;  
}
```

# Common sub-expression elimination

We can cache the value of  $v * 3$ .

Assuming  $v$  doesn't change

- depends on data flow analysis!

```
x := foo(v * 3);
```

```
if (v * 3 < 9) {  
    ...  
}
```



```
t := v * 3;  
x := foo(t);
```

```
if (t < 9) {  
    ...  
}
```



# Where is the pattern?

We have seen some optimizations.

- All seem ad-hoc.
- All depend on knowing something about their **context**...

- Some optimizations are local.
- Some optimizations act over multiple basic blocks.


Most optimizing compilers now use SSA-form to unify many optimizations.



# DU, UD chains

- DU chain = “definition use” chain
  - directed arc(s) from each variable definition to the use(s) of that variable
- UD chain = “use definition”
  - directed arc(s) from a variable use to the instruction defining that variable
- Both are implemented as graphs
  - common technique before SSA

# Example: DU, UD chains



# Static Single-Assignment

- Invariant on instruction stream
  - Every virtual register has one (static) definition site
  - Never re-assign a virtual register.

This is straightforward for straight-line code.

```
a ← x * y
b ← a - 1
a ← y * b
b ← x * 4
a ← a + b
```

```
a1 ← x * y
b1 ← a1 - 1
a2 ← y * b1
b2 ← x * 4
a3 ← a2 + b2
```



## SSA (2)

```
a ← x * y
b ← a - 1
a ← y * b
b ← x * 4
a ← a + b
```

```
a1 ← x * y
b1 ← a1 - 1
a2 ← y * b1
b2 ← x * 4
a3 ← a2 + b2
```

- $a_1, a_2, a_3$  are distinct virtual registers.
- They *may* map to the same physical register.

# SSA (3)


```
a ← x * y
b ← a - 1
a ← y * b
b ← x * 4
a ← a + b
```

```
a1 ← x * y
b1 ← a1 - 1
a2 ← y * b1
b2 ← x * 4
a3 ← a2 + b2
```

- We now **know** the value of  $a_1$  **for all time**.
- It is **never** reassigned.
  - ➔ constant folding, etc become easier.
  - if ( $a_1 < 10$ ) { ... }





# SSA Control Flow (1)



What about control flow?


- How do we give SSA definitions for  $a$  and  $c$ ?

# SSA Control Flow (2)





Block 1 is straightforward...

# SSA Control Flow (3)



In block 2, we define a new  $a$ .  
Which ' $a$ ' do we use in block 3?


# SSA Control Flow (4)



Use a  $\Phi$ -node...


# The $\Phi$ -node

$\Phi$ -nodes: means something like a “choice” between values  $a_N$  and  $a_M$



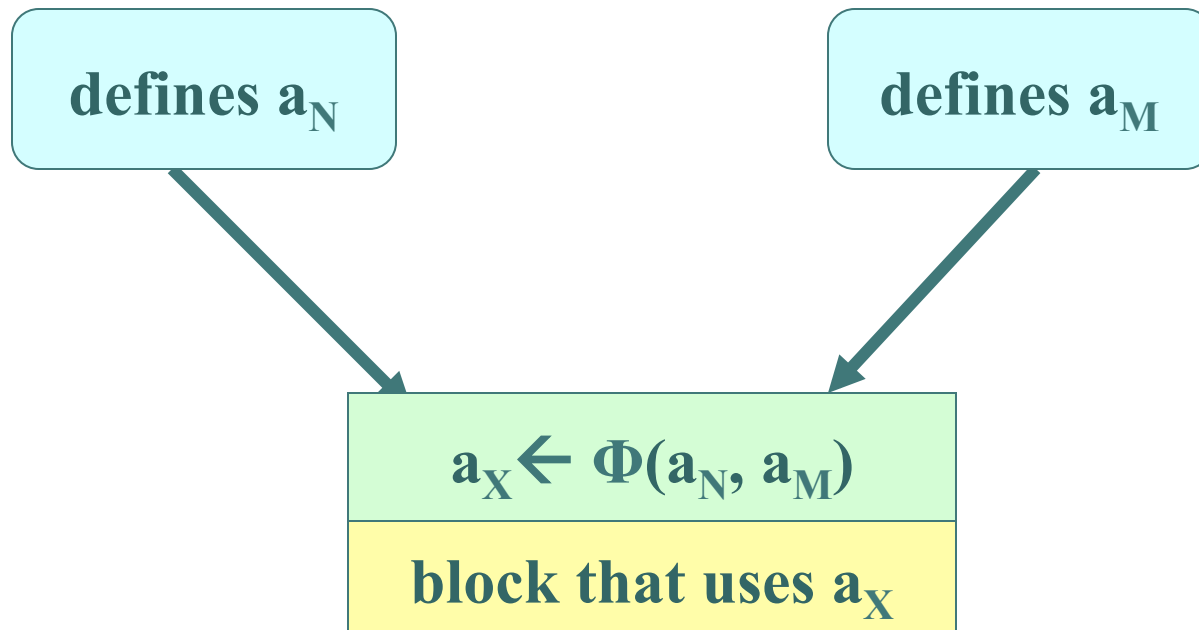
# The $\Phi$ -node

$\Phi$ -nodes: means something like a “choice” between values  $a_N$  and  $a_M$



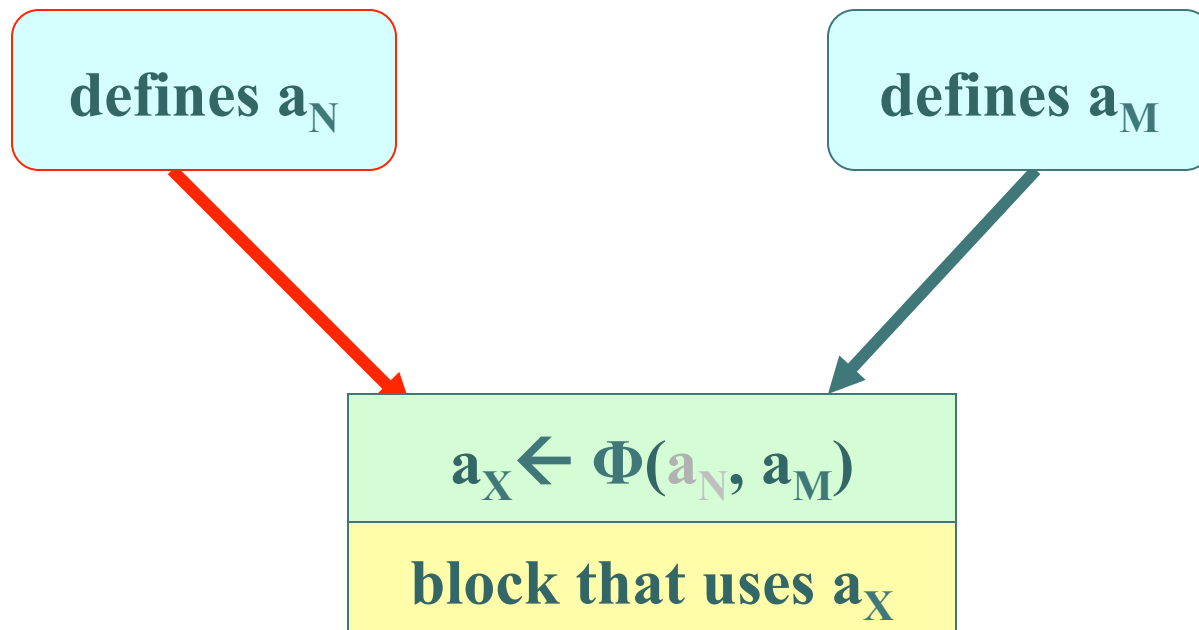
# The $\Phi$ -node

SSA is one of the most important new idioms for compiler construction in the last 20 years!



# The $\Phi$ -node

If you come from the left-hand basic block, the  $\Phi$  operation is a copy from  $a_N$  to  $a_X$








# The $\Phi$ -node

Think of the  $\Phi$  operation as a “magic move” operation, that copies one of its operands to its destination.

$$a_X \leftarrow \Phi(a_N, a_M)$$


block that uses  $a_X$

# SSA Control Flow (5)



What about instruction generation?

# SSA Control Flow (6)



```
r1 ← M[x]
r2 ← 0
if (r1 < 4) goto L1:
goto L2:
```

L1:

```
r3 ← r1
goto L3:
```

L2:

```
goto L3:
```

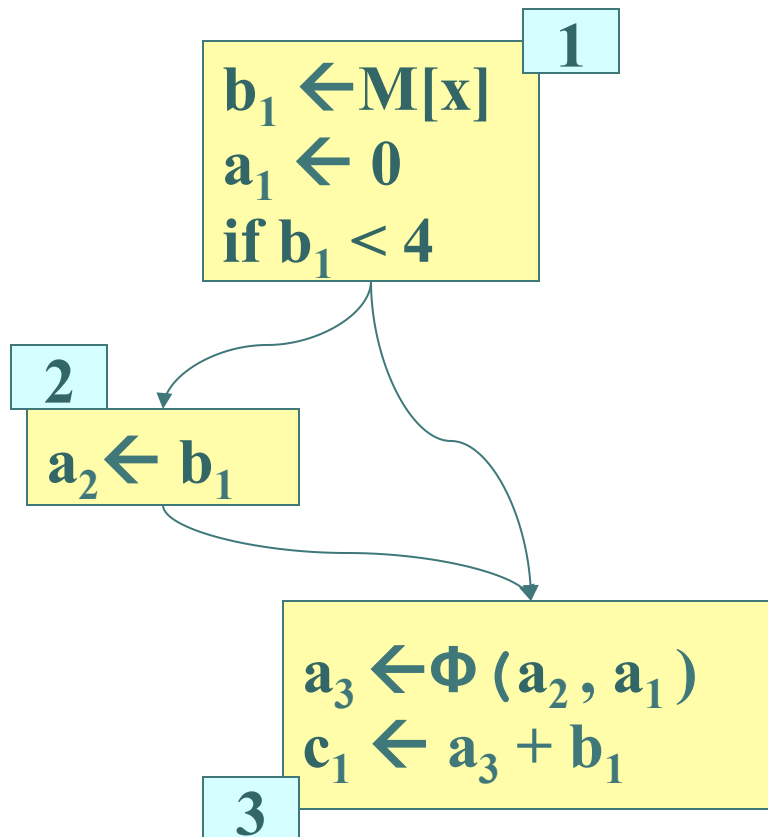
L3:

```
r5 ← ?? + r1
```

may be  
r2 or r3

What about instruction generation?

# SSA Control Flow (7)



Insert extra moves...

```
r1 ← M[x]
r2 ← 0
if (r1 < 4) goto L1:
goto L2:
```

L1:

```
r3 ← r1
r4 ← r3
goto L3:
```

L2:

```
r4 ← r2
goto L3:
```


L3:

```
r5 ← r4 + r1
```

**Φ implemented via moves.**

Note: r4 is new

# SSA Control Flow (8)



```
r1 ← M[x]
r2 ← 0
if (r1 < 4) goto L1:
goto L2:
```

L1:

```
r2 ← r1
r2 ← r2
```

---

```
goto L3:
```

L2:

```
r2 ← r2
```

---

```
goto L3:
```

L3:

```
r5 ← r2 + r1
```


Best case: if all 'a' s map to r2 (depends on register allocation)



# $\Phi$ Summary

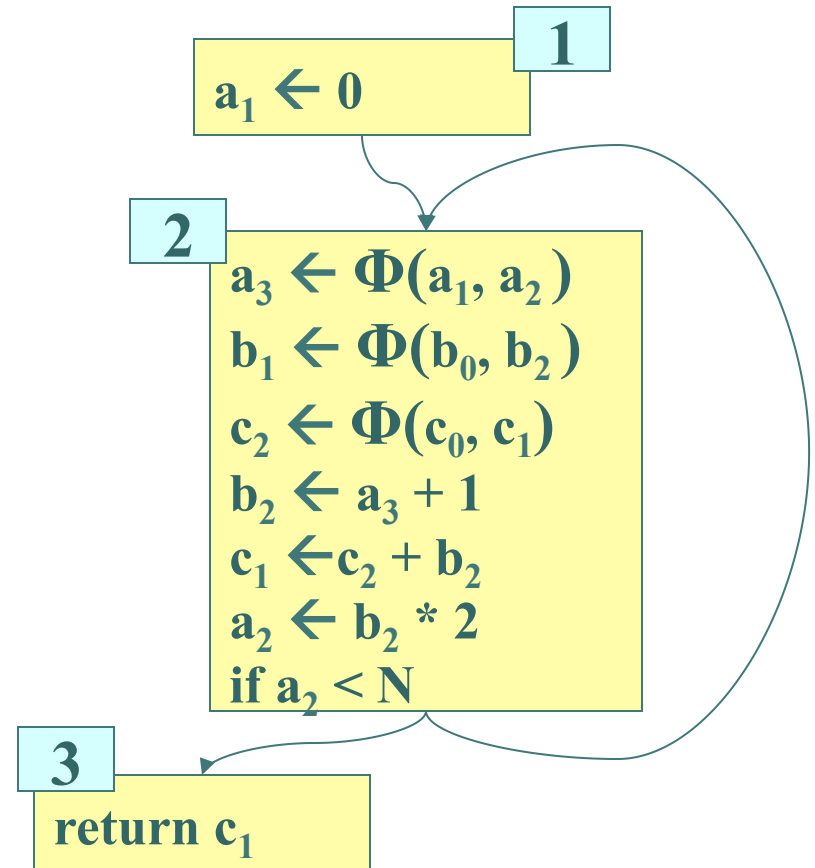

- For optimization purposes, assume  $\Phi$  instructions are just normal instructions
  - e.g.,  $a_3 \leftarrow \Phi(a_2, a_1)$  is like an ordinary instruction for the purposes of data flow
- Typically they can be removed at register allocation time.
  - e.g., if  $a_2, a_1$  are placed in the same register, then the above node may be removed
- If not, extra moves need to be inserted in previous basic blocks to implement  $\Phi$  node.

# Loops and $\Phi$ nodes



This is translated just like the straight line example...

# Loops and $\Phi$ nodes



This is translated just like the straight line example...





# Is this really Static Single Assignment?

Each time round a loop, the same virtual register is re-assigned.

- But it is defined **statically** only once.
- Multiple **dynamic** definitions are allowed.



# Using SSA: dead code elimination

Consider

**$v1 \leftarrow x3 + y2$**

$v1$  has no uses.

- easy to tell by inspecting SSA code
- This must be the only definition of  $v1$ .
- So  $v1$  is dead.
- This assignment to  $v1$  can be removed!

- The operation must not have side-effects.
  - i.e., it can't be print, writing memory, etc.
- Removing this assignment can cause other instruction to also become dead.
  - i.e., deletes uses of  $x, y$

Happens in practice inside optimizing compilers.



# Using SSA: Constant propagation

- Descend through the instructions, remembering constants and properties.

$x \leftarrow 9$

...

```
while(...) {  
     $y \leftarrow x + 2$ ;  
}
```

$x \leftarrow \text{const}$

$\{ x \rightarrow \langle \text{const} \rangle \}$

$x \leftarrow a \text{ op } b$

if a and b are  
constants, then evaluate:

$\{ x \rightarrow \langle \text{const} \rangle \}$

**if ( $x > \text{const}$ )**

- on one branch  $x > \text{const}$
- on the other branch  $x \leq \text{const}$