

Compilers I

Introduction to Compiler Optimization

Dr. William L. Harrison

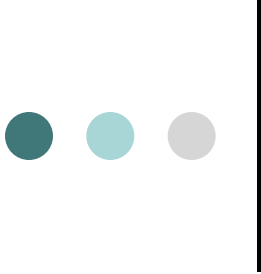
harrisonwl@missouri.edu



Today's Class

Today: Optimizations

- Basic Optimizations
- Static Single-Assignment



What does an optimizing compiler do?



Optimizing compiler attempts to:

- ♦ Eliminate language abstractions
- ♦ Map source program to target machine efficiently.
 - ♦ Use the hardware well!
- ♦ Equal the efficiency of a good assembly programmer.



Dimensions of Optimization

You can optimize for a number of things.

- machine cycles (speed)
- binary size
- heap usage
- to use special hardware features
 - parallel computing

Minimizing runtime cycles is the typical measure to optimize.



Proebsting's Law

- Moore's Law: Computing power doubles every 18 months.
 - This has been true for 40 years!
- Proebsting's Law: **Compiler Advances** double computing power every 18 *Years*
 - Better optimizations
 - Better use of new hardware features



Optimizing compilers allow efficient abstractions

- We can engineer our software applications using abstractions like OO-style objects, etc.
 - Ease of maintenance.
 - Almost all design patterns are abstractions.
- The compiler can remove much of the overheads associated with using these abstractions.



Appel's Compiler Theorem

Theorem: For any optimizing compiler, there exists a better one.

“full employment theorem for compiler writers...”

Appel's point was you can always add a new optimizations to make any compiler better.

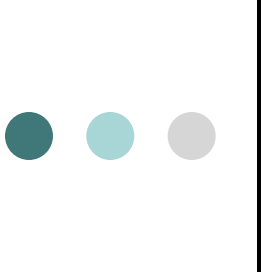
What is an optimization?

- The output program must have the same **observable behavior**, as defined in the semantics of the source language.

```
int foo(int i) {  
    int j = 9;  
    return i + j;  
}
```

≈

```
int foo(int i) {  
    return i+9;  
}
```

Optimizations can make your working program fail!

- The output program must have the same **observable behavior**, as defined in the semantics of the source language.
- If a program steps outside the semantically defined part of the language, all bets are off.
 - Example: using an un-initialized variable in C.
The value might be different between optimized and non-optimized.
- Your program was faulty to start with!
 - It just happened to work...



Optimizations

- Optimizations change the program internally
- This enables more optimizations...
 - Some optimizations are ***enabling*** optimizations
 - Some optimizations are ***true*** optimizations



Optimization Camps

Optimizations that move computation from runtime to compile time.

- *Examples: #ifdef, asserts, constant folding*

Optimizations that replace some instructions with less expensive instructions.

- *Examples: peephole optimizations*

Optimizations that move a computation to a less expensive place.

- *Examples: loop invariant hoisting*


Compile Time Evaluation

Sometimes
constants can be
evaluated at
compile time.

Caveats

x must be constant.

```
final int x = 9;  
if (x < 10) {  
    ...  
}
```



```
final int x = 9;  
if (true) {  
    ...  
}
```




Constant folding

- Constant folding is a true optimization
 - The work that would have been done at runtime is no longer needed.
- Constant folding is **also** an enabling optimization.
 - On previous page, we can eliminate the if/then test.
 - Constant folding can give rise to **dead-code**.

Peephole Optimizations

Consider these instructions

| | |
|------|----------------------------------|
| ADDI | $r_1 \leftarrow V_a$ |
| ADD | $r_2 \leftarrow \text{fp} + r_1$ |
| LOAD | $r_3 \leftarrow M[r_2 + 0]$ |



But

- V_a is a constant.
- fp is a register.

→ so we could write

| |
|--|
| LOAD $r_3 \leftarrow M[\text{fp} + V_a]$ |
|--|

Works inside basic blocks

Need to know about
usage of r_1 and r_2 .

Procedure inlining

Two steps to inlining

- Substitute procedure body at call site
- Watch for name clashes
 - Fix variable names if clash

```
int foo(int x) {  
    return x * x;  
}
```

```
...  
    v := foo(y);  
...
```

Saves cost of
function call



```
...  
    v := y * y;  
...
```

Loop invariants

Assuming v does not change inside this loop

Can “lift” the computation of $v * 2$ out of the loop.

```
v := ...  
...  
while (x < 10)  
{  
    y = v * 2;  
    x = x + y;  
}
```



It is almost
always good
to move things
out of loops.

```
v := ...  
...  
y := v * 2;  
while (x < 10) {  
    x = x + y;  
}
```


Common sub-expression elimination

We can cache the value of $v * 3$.

Assuming v doesn't change

- depends on data flow analysis!

```
x := foo(v * 3);
```

```
if (v * 3 < 9) {  
    ...  
}
```



```
t := v * 3;  
x := foo(t);
```

```
if (t < 9) {  
    ...  
}
```



Where is the pattern?

We have seen some optimizations.

- All seem ad-hoc.
- All depend on knowing something about their **context**...

- Some optimizations are local.
- Some optimizations act over multiple basic blocks.

Most optimizing compilers now use SSA-form to unify many optimizations.



A survey of optimizations



Optimizing compilers

There's more to performance than asymptotic complexity

- Constant factors matter too!
 - Easily see 10:1 performance range depending on how code is written
 - Must optimize at multiple levels:
 - algorithm, data representations, procedures, and loops
- Understand target system to optimize performance
 - How programs are compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality



Optimizing compilers

- Provide efficient mapping of program to machine
 - register allocation
 - code selection and ordering (scheduling)
 - dead code elimination
 - eliminating minor inefficiencies
- Don't (usually) improve asymptotic efficiency
 - up to programmer to select best overall algorithm
 - big-O savings are (often) more important than constant factors (but constant factors can make a big difference)
- Have difficulty overcoming “optimization blockers”
 - potential memory aliasing
 - potential procedure side-effects



Limitations of optimizing compilers

○ Fundamental Constraint

- Must not cause any change in observable program behavior under any possible condition
 - Can prevent making optimizations when would only affect behavior under pathological conditions.

○ Behavior that may be obvious to the programmer can be obfuscated by languages and coding styles

- e.g., Data ranges may be more limited than variable types suggest

○ Most analysis is performed only within procedures

- Whole-program analysis is too expensive in most cases

○ Most analysis is based only on *static* information

- Compiler has difficulty anticipating run-time inputs

When in doubt, the compiler must be conservative



What to optimize?

You can optimize for a number of things:

- machine cycles (speed)
- binary size
- heap usage
- to use special hardware features
 - parallel computing
- compile time (?!)

What to optimize?

You can optimize for a number of things:

- machine cycles (speed)

- binary size

- heap usage

- to use special hardware features

 - parallel computing

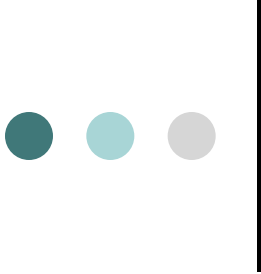
- compile time (?!)

(what we usually
care about most)



Proebsting's Law

- Moore's Law: computing power (transistor density) doubles every 18 months.
- Proebsting's Law: advances in compiler optimizations double computing power every 18 *years*.



Optimization: a transformation
that (hopefully) improves
performance


The output program must have the same
observable behavior, as defined in the
semantics of the source language.

```
int foo(int i) {  
    int j = 9;  
    return i + j;  
}
```

≈


```
int foo(int i) {  
    return i + 9;  
}
```

Middle end optimizations




Optimizations in the middle end are backend-agnostic and usually operate on the CFG, taking advantage of various dataflow analyses.

Middle end optimizations



- Constant folding/propagation
- Copy propagation
- Algebraic simplification
- Common subexpression elimination
- Dead code elimination
- Loop optimizations
- Function inlining

Middle end optimizations



- Constant folding/propagation
- Copy propagation
- Algebraic simplification
- Common subexpression elimination
- Dead code elimination
- Loop optimizations
- Function inlining

Constant folding

- Evaluate constant expressions at compile time
- Only possible when side-effect freeness guaranteed

```
c ← 1 + 3
```



```
c ← 4
```

```
not true
```



```
false
```

Constant propagation

Variables that have constant value, e.g. $c := 3$

- Later uses of c can be replaced by the constant
- If no change of c between!

```
b ← 3  
c ← 1 + b  
d ← b + c
```



```
b ← 3  
c ← 1 + 3  
d ← 3 + c
```

Analysis needed, as b can be assigned more than once!

Constant propagation

Variables that have constant value, e.g. $c := 3$

- Later uses of c can be replaced by the constant
- If no change of c between!

Ready for constant folding now.

```
b ← 3  
c ← 1 + b  
d ← b + c
```



```
b ← 3  
c ← 1 + 3  
d ← 3 + c
```

Analysis needed, as b can be assigned more than once!

Copy propagation

- for a statement $x \leftarrow y$
- replace later uses of x with y , if x and y have not been changed.

```
x ← y  
c ← 1 + x  
d ← x + c
```



```
x ← y  
c ← 1 + y  
d ← y + c
```

Analysis needed, as y and x can be assigned more than once!

Algebraic simplification

Use algebraic properties to simplify expressions

`-(-i)`



`i`

`b or true`



`true`

Common subexpression elimination

- “Cache” an expression to avoid re-evaluation.
- Only works if the registers used in the expression remain unchanged between uses!

```
x = foo(v * 3);
```

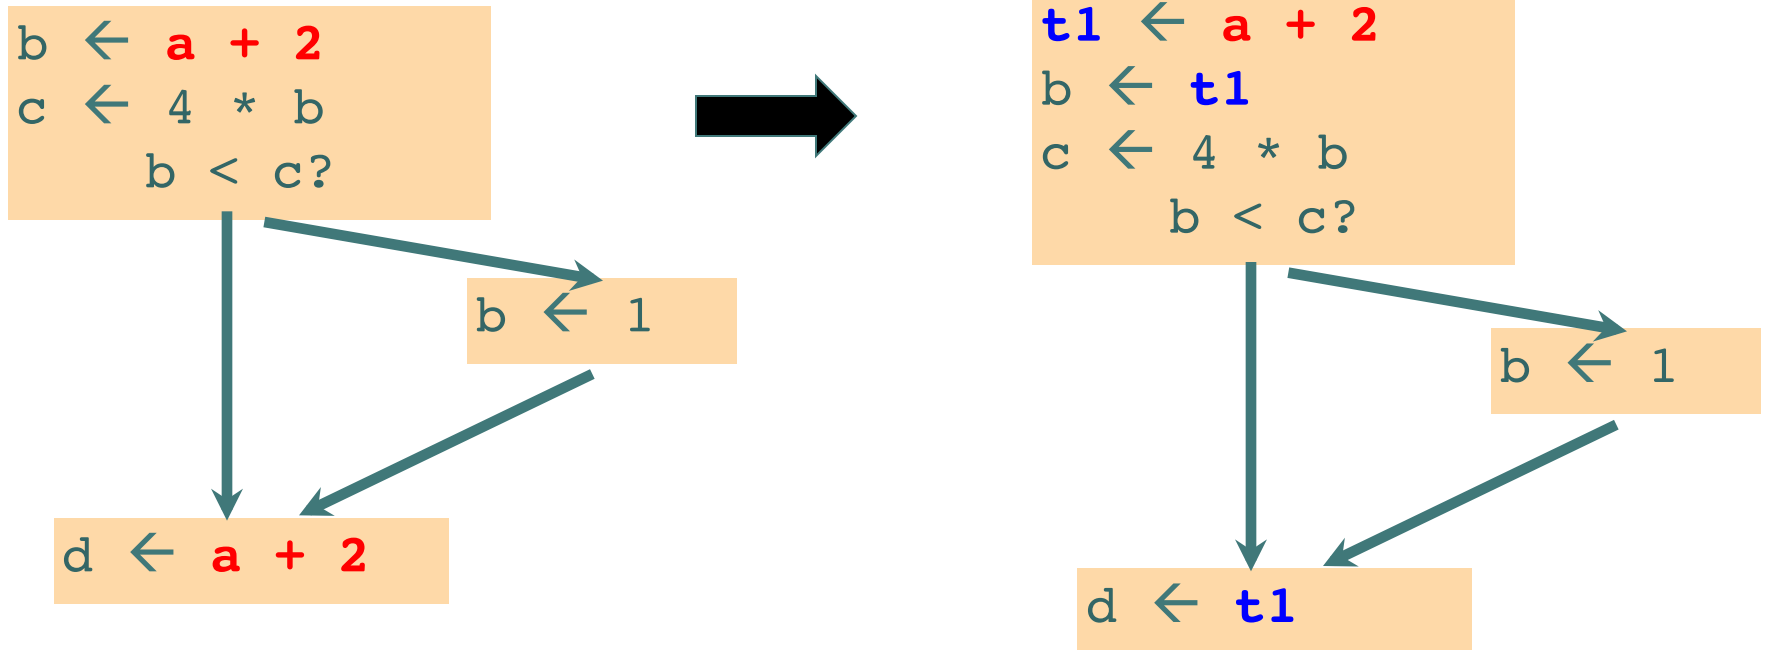
```
if (v * 3 < 9) {  
    ...  
}
```



```
t = v * 3;  
x = foo(t);
```


```
if (t < 9) {  
    ...  
}
```

Common subexpression elimination



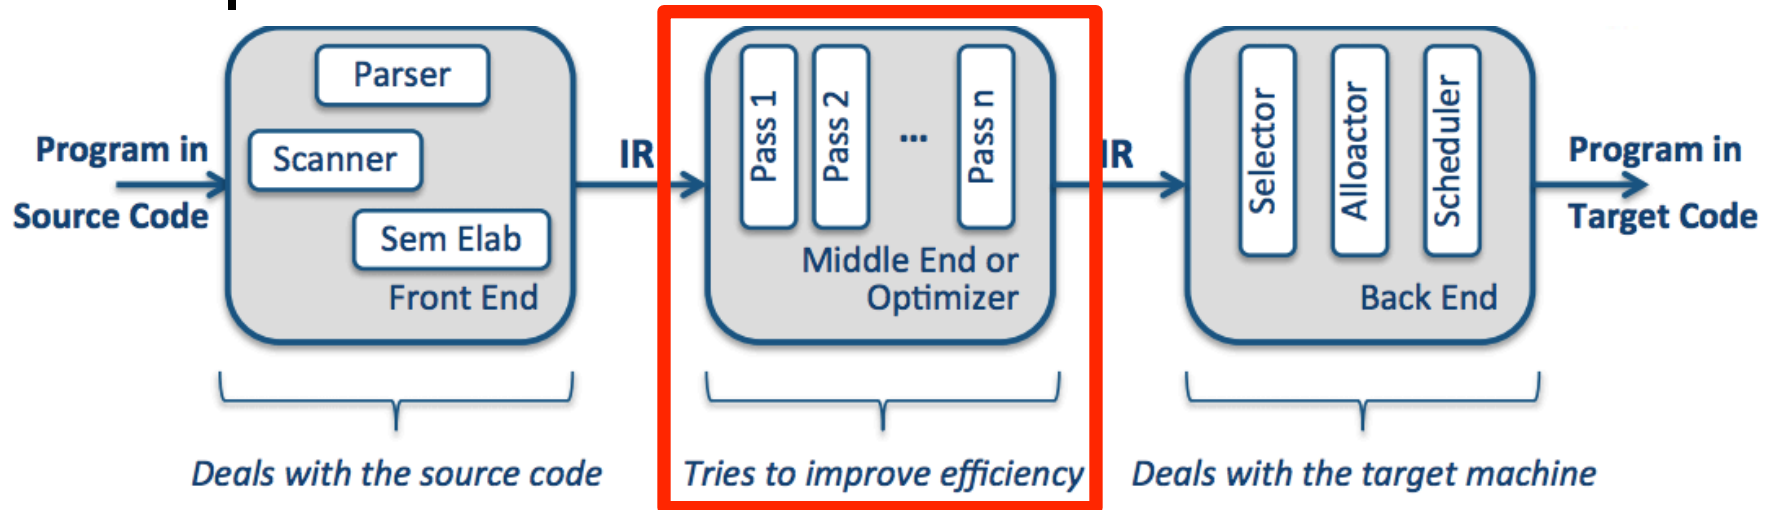
* The Moral: inter-block CSE trickier than intra-block CSE

Middle end optimizations



- Constant folding/propagation
- Copy propagation
- Algebraic simplification
- Common subexpression elimination
- Dead code elimination
- Loop optimizations
- Function inlining

Middle end optimizations



- Constant folding/propagation
- Copy propagation
- Algebraic simplification
- Common subexpression elimination
- Dead code elimination
- Loop optimizations
- Function inlining

Dead code elimination

- Remove unnecessary code
 - e.g. variables assigned but never read:

```
b ← 3  
c ← 1 + 3  
d ← 3 + c
```



```
c ← 1 + 3  
d ← 3 + c
```

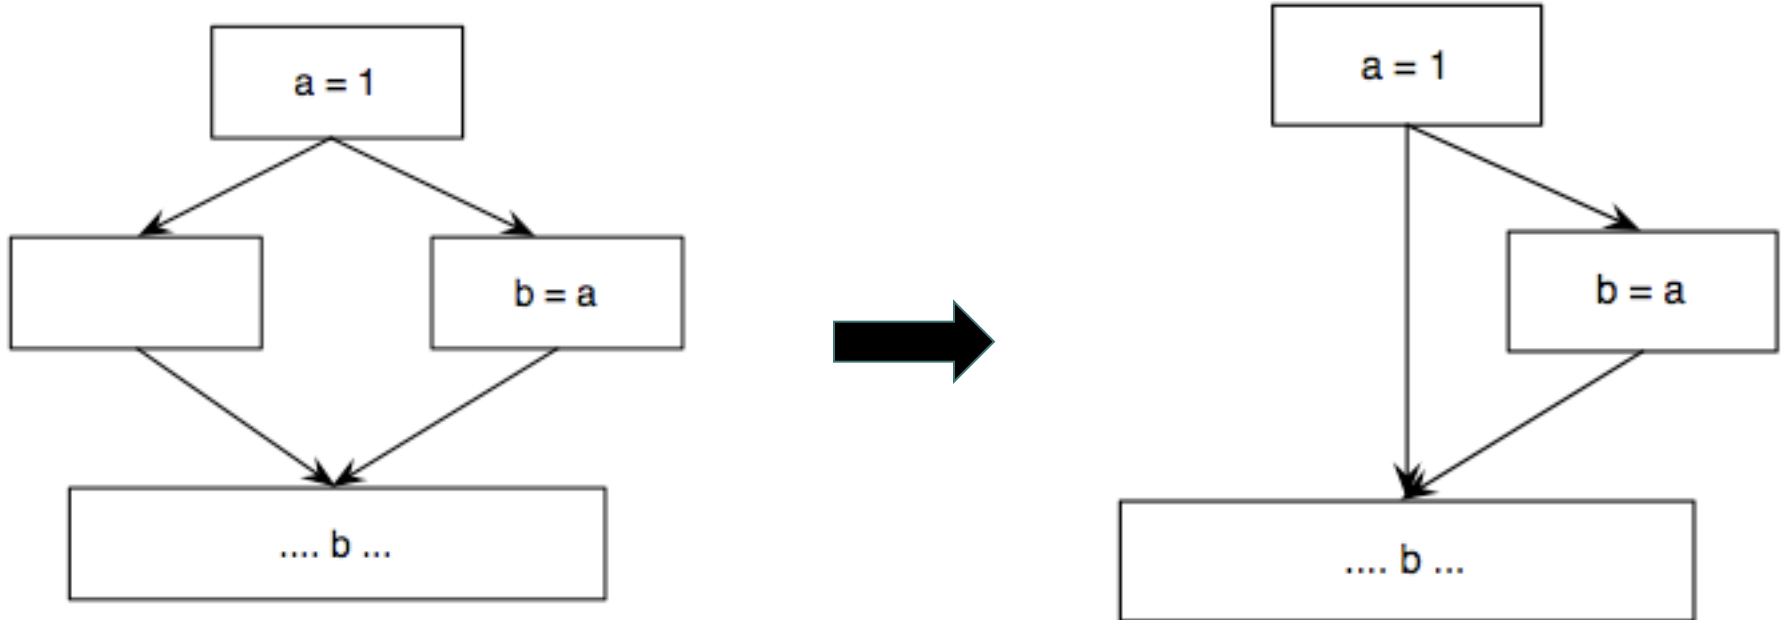
– blocks never reached:

```
if (false) {  
    a ← 5  
}
```




```
if (false) {}
```

Dead code elimination



Other CFG simplifications: "Fusion" of basic blocks




AKA, "Code Straightening"



DU, UD chains

- DU chain = “definition use” chain
 - directed arc(s) from each variable definition to the use(s) of that variable
- UD chain = “use definition”
 - directed arc(s) from a variable use to the instruction defining that variable
- Both are implemented as graphs
 - common technique before SSA

Example: DU, UD chains



Static Single-Assignment

- Invariant on instruction stream
 - Every virtual register has one (static) definition site
 - Never re-assign a virtual register.

This is straightforward for straight-line code.

```
a ← x * y
b ← a - 1
a ← y * b
b ← x * 4
a ← a + b
```

```
a1 ← x * y
b1 ← a1 - 1
a2 ← y * b1
b2 ← x * 4
a3 ← a2 + b2
```



SSA (2)

```
a ← x * y
b ← a - 1
a ← y * b
b ← x * 4
a ← a + b
```

```
a1 ← x * y
b1 ← a1 - 1
a2 ← y * b1
b2 ← x * 4
a3 ← a2 + b2
```

- a_1 , a_2 , a_3 are distinct virtual registers.
- They *may* map to the same physical register.


SSA (3)

```
a ← x * y
b ← a - 1
a ← y * b
b ← x * 4
a ← a + b
```

```
a1 ← x * y
b1 ← a1 - 1
a2 ← y * b1
b2 ← x * 4
a3 ← a2 + b2
```

- We now **know** the value of a_1 **for all time**.
- It is **never** reassigned.
 - ➔ constant folding, etc become easier.
 - if ($a_1 < 10$) { ... }



SSA Control Flow (1)



What about control flow?

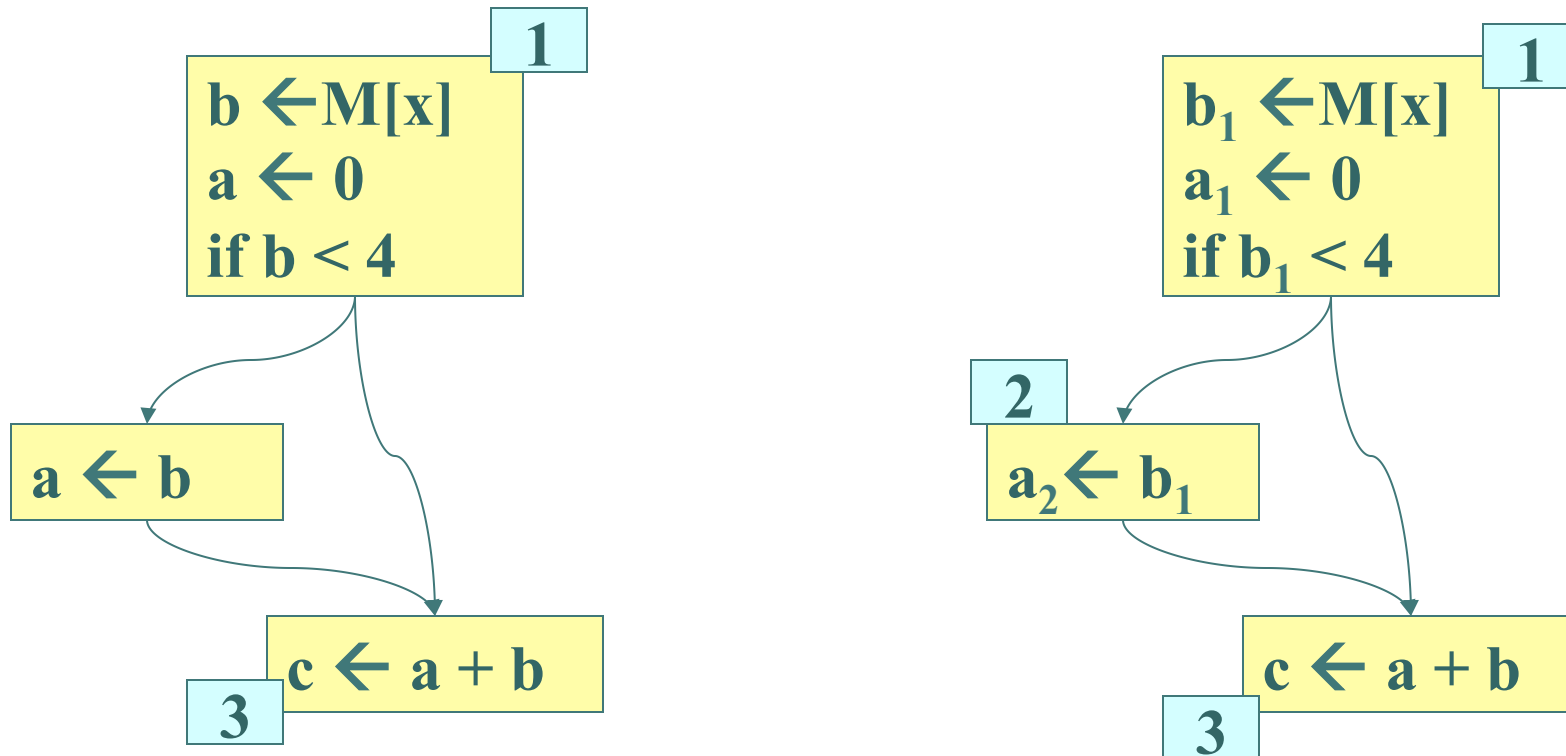
- How do we give SSA definitions for a and c ?

SSA Control Flow (2)





Block 1 is straightforward...

SSA Control Flow (3)



In block 2, we define a new a .
Which ' a ' do we use in block 3?


SSA Control Flow (4)



Use a Φ -node...

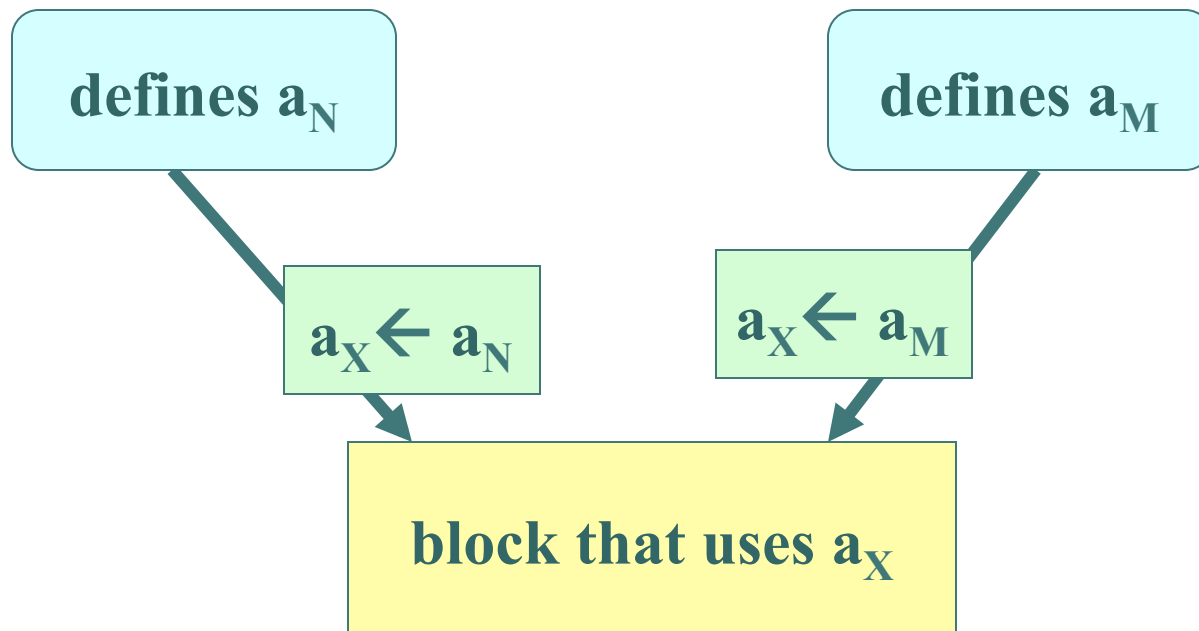
Φ-node?

Φ-nodes: means something like a “choice” between values a_N and a_M



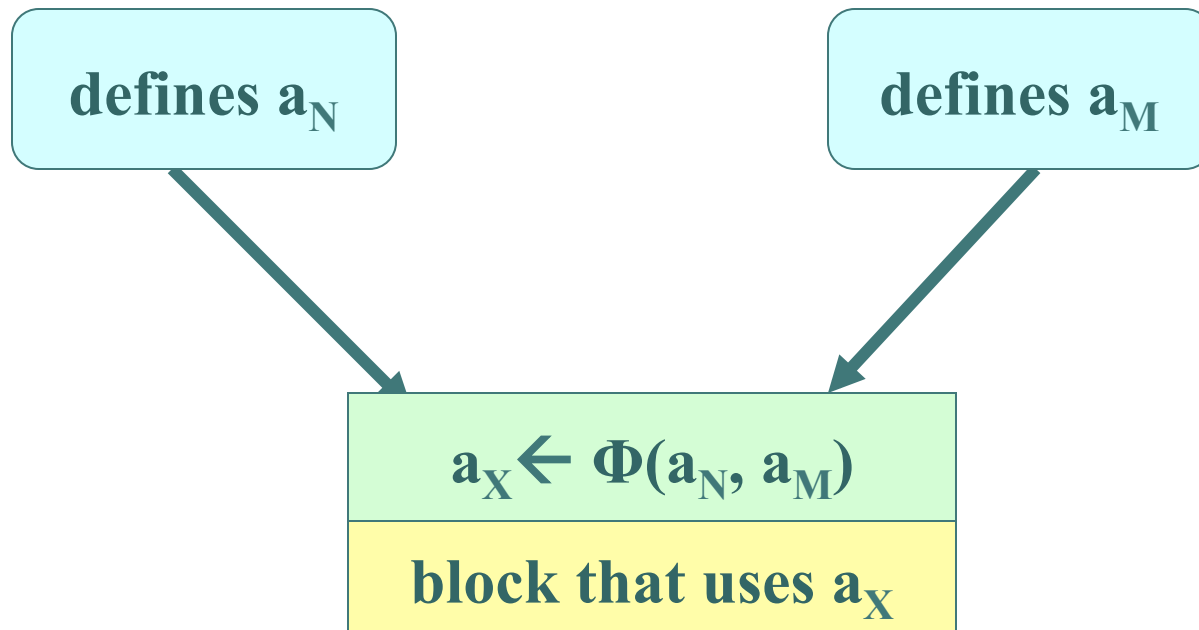
The Φ -node

Φ -nodes: means something like a “choice” between values a_N and a_M



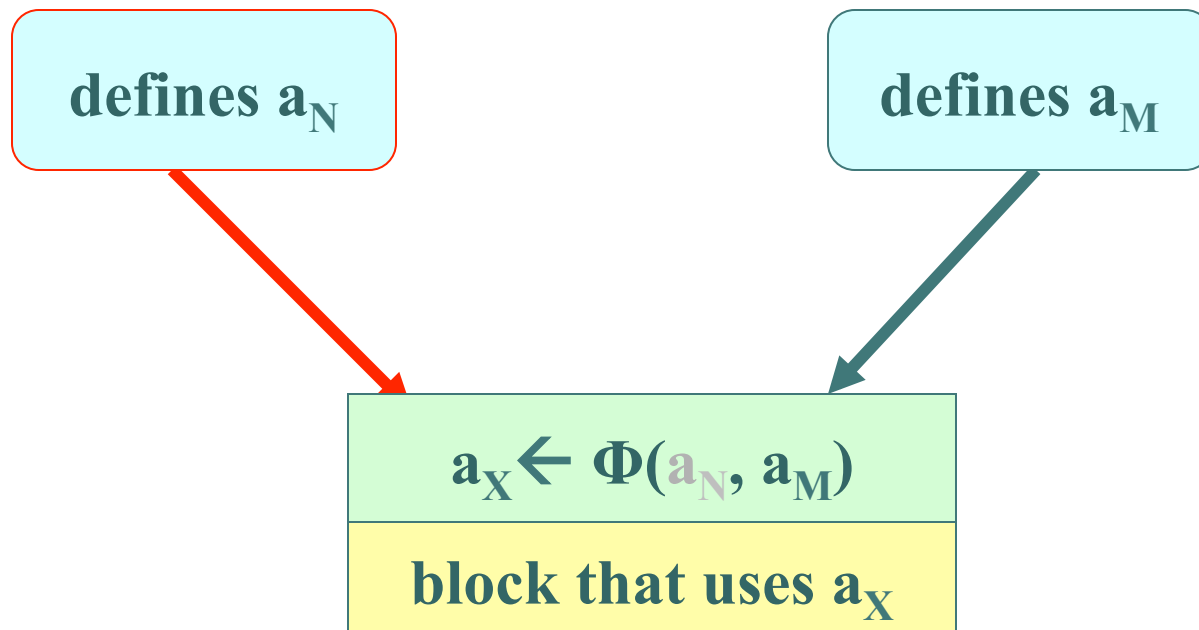
The Φ -node

SSA is one of the most important new idioms for compiler construction in the last 20 years!



The Φ -node

If you come from the left-hand basic block, the Φ operation is a copy from a_N to a_X





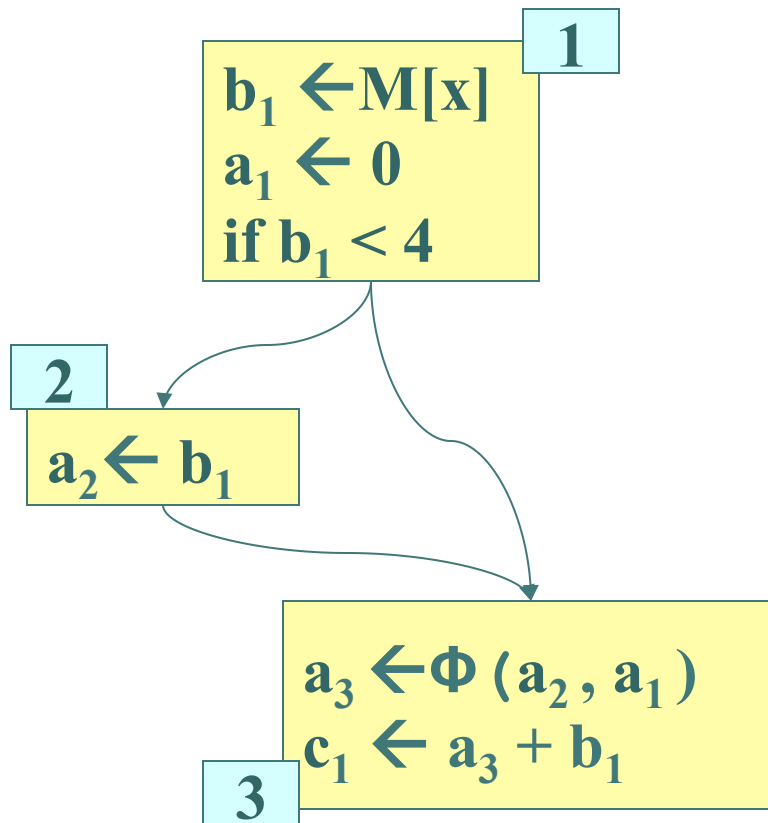
The Φ -node

Think of the Φ operation as a “magic move” operation, that copies one of its operands to its destination.

$$a_X \leftarrow \Phi(a_N, a_M)$$


block that uses a_X

SSA Control Flow (5)



What about instruction generation?

SSA Control Flow (6)



```
r1 ← M[x]
r2 ← 0
if (r1 < 4) goto L1:
goto L2:
```

L1:

```
r3 ← r1
goto L3:
```

L2:

```
goto L3:
```


L3:

```
r5 ← ?? + r1
```

may be
r2 or r3

What about instruction generation?

SSA Control Flow (7)



Insert extra moves...

```
r1 ← M[x]
r2 ← 0
if (r1 < 4) goto L1:
goto L2:
```

L1:

```
r3 ← r1
r4 ← r3
goto L3:
```

L2:

```
r4 ← r2
goto L3:
```

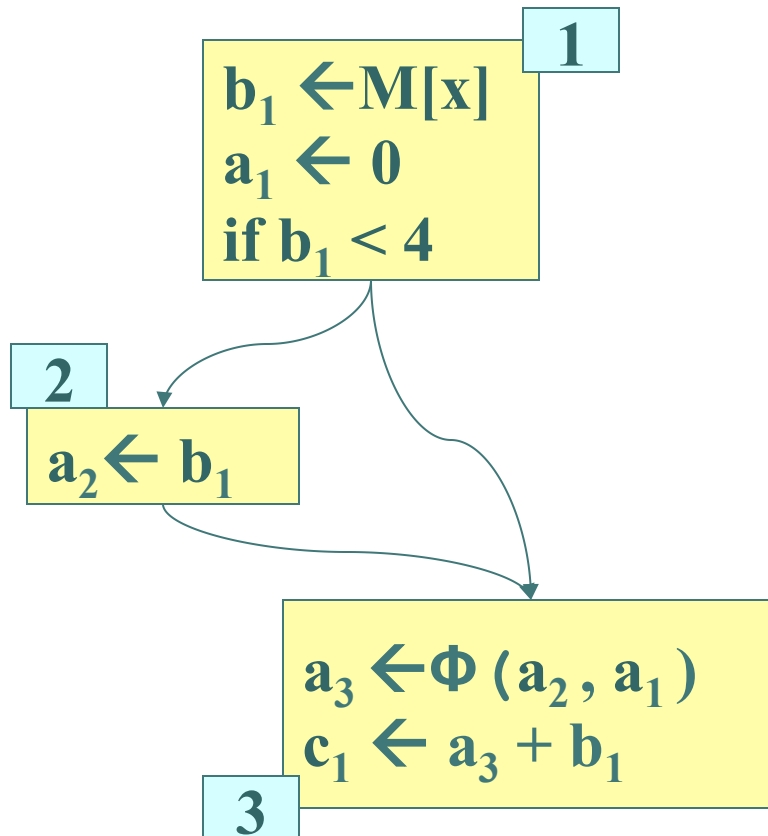
L3:

```
r5 ← r4 + r1
```

Φ implemented
via moves.

Note: r4 is new

SSA Control Flow (8)



```
r1 ← M[x]
r2 ← 0
if (r1 < 4) goto L1:
goto L2:
```

L1:

```
r2 ← r1
r2 ← r2
```

```
goto L3:
```

L2:

```
r2 ← r2
```

```
goto L3:
```

L3:

```
r5 ← r2 + r1
```


Best case: if all 'a' s map to r2 (depends on register allocation)



Φ Summary

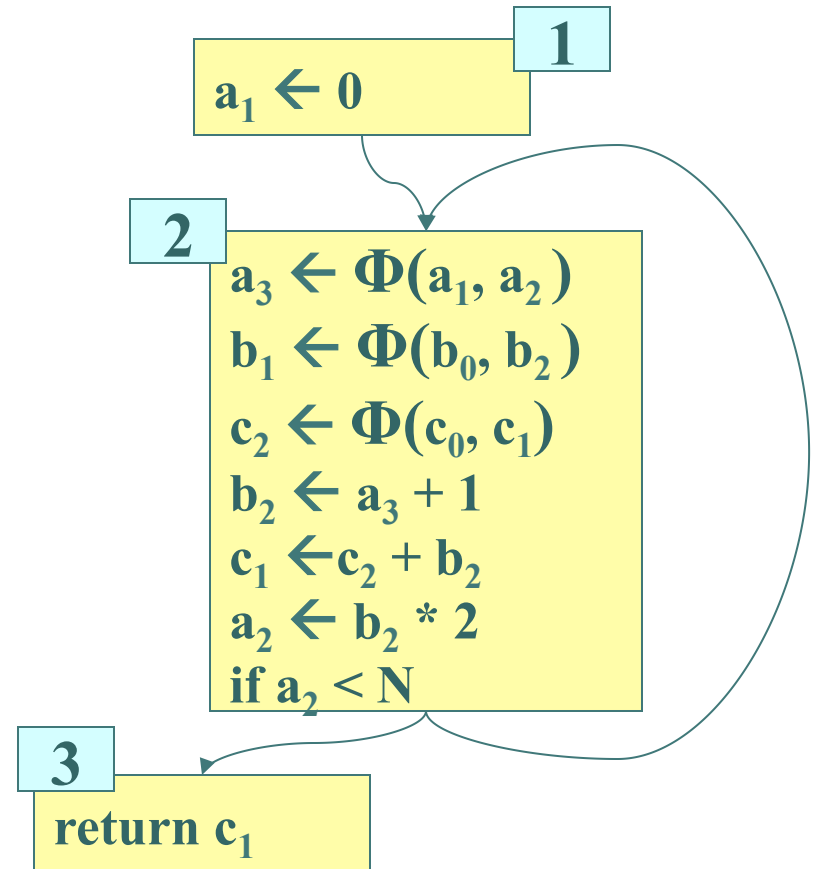

- For optimization purposes, assume Φ instructions are just normal instructions
 - e.g., $a_3 \leftarrow \Phi(a_2, a_1)$ is like an ordinary instruction for the purposes of data flow
- Typically they can be removed at register allocation time.
 - e.g., if a_2, a_1 are placed in the same register, then the above node may be removed
- If not, extra moves need to be inserted in previous basic blocks to implement Φ node.

Loops and Φ nodes



This is translated just like the straight line example...

Loops and Φ nodes



This is translated just like the straight line example...



Is this really Static Single Assignment?

Each time round a loop, the same virtual register is re-assigned.

- But it is defined **statically** only once.
- Multiple **dynamic** definitions are allowed.



Using SSA: dead code elimination

Consider

$v1 \leftarrow x3 + y2$

If $v1$ has no uses, then

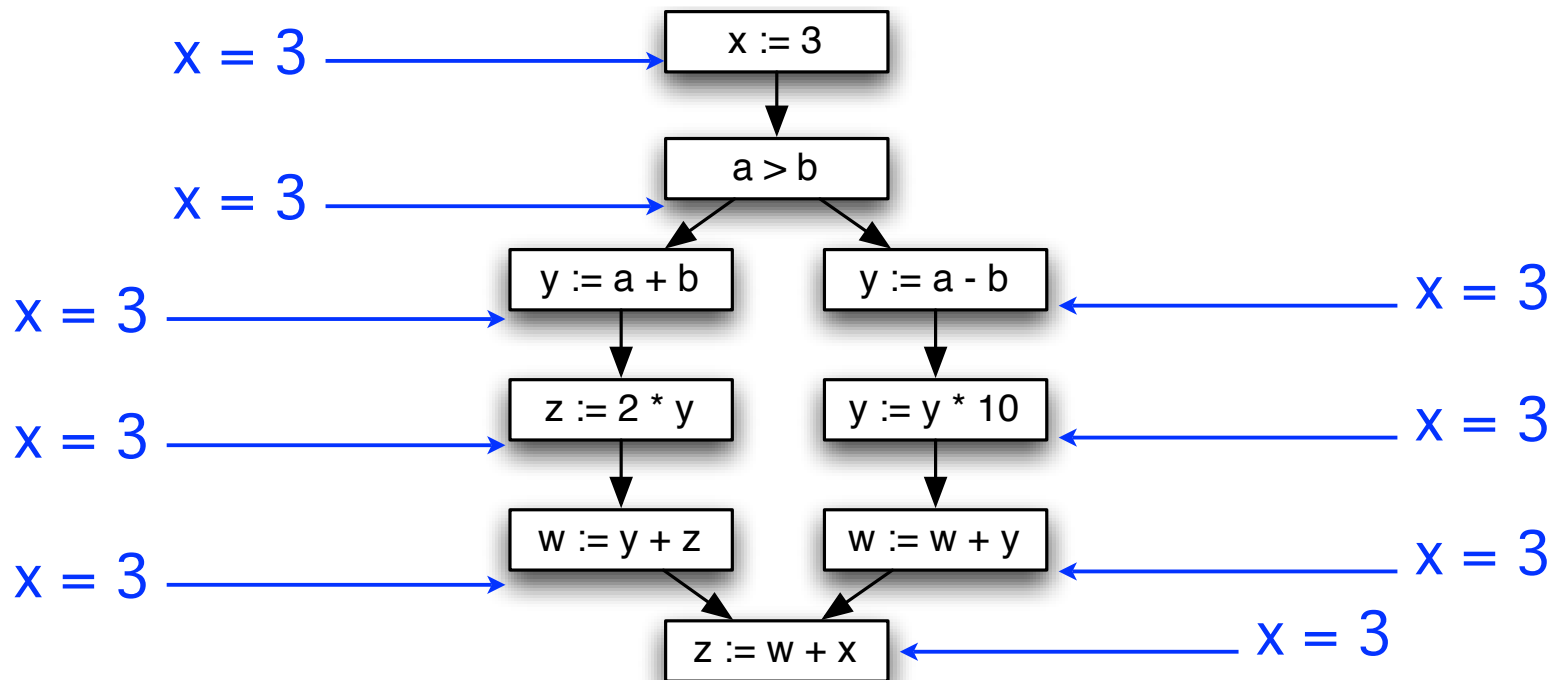
- easy to tell by inspecting SSA code
- This must be the only definition of $v1$.
- So $v1$ is dead.
- This assignment to $v1$ can be removed!

- The operation must not have side-effects.
 - i.e., it can't be print, writing memory, etc.
- Removing this assignment can cause other instruction to also become dead.
 - i.e., deletes uses of x, y

Happens in practice inside optimizing compilers.

Using SSA: Constant propagation

Before SSA



Using SSA: Constant propagation

After SSA Why is this a big win over the previous slide?

