



Compilers I

Loop Optimizations

Dr William Harrison



Loop Optimizations

- Empirical studies show that much of program execution occurs within loops
 - So, it makes sense to identify loops & concentrate optimization efforts on that code
 - But what is a loop?
 - Within source code, it's obvious
 - Within IR/CFG, however,...
 - Loops aren't precisely what your first guess would be



Simple loop optimization

within the source code, that is...

`for (i = 1; i++ ; i = 10) { c }` **→** `c ; ... ; c`

10 times

Caveat:

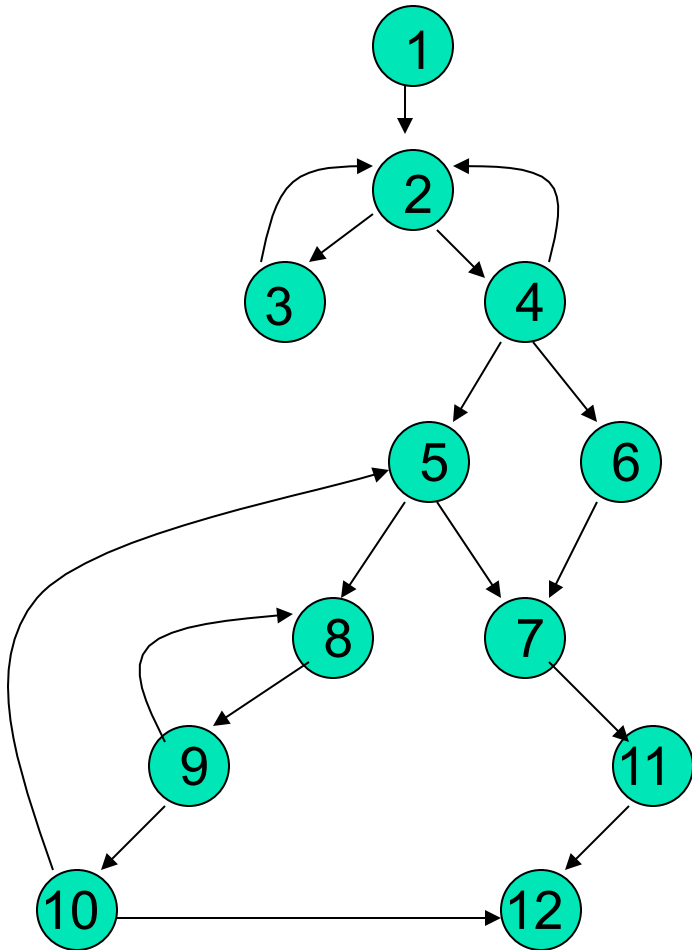
- seems reasonable (removes branches)
- not terribly general: would like to use on while loops but how do you identify the induction variable “i”?
- also, how does this apply at the IR/CFG level?



Ex: Loop invariant code hoisting

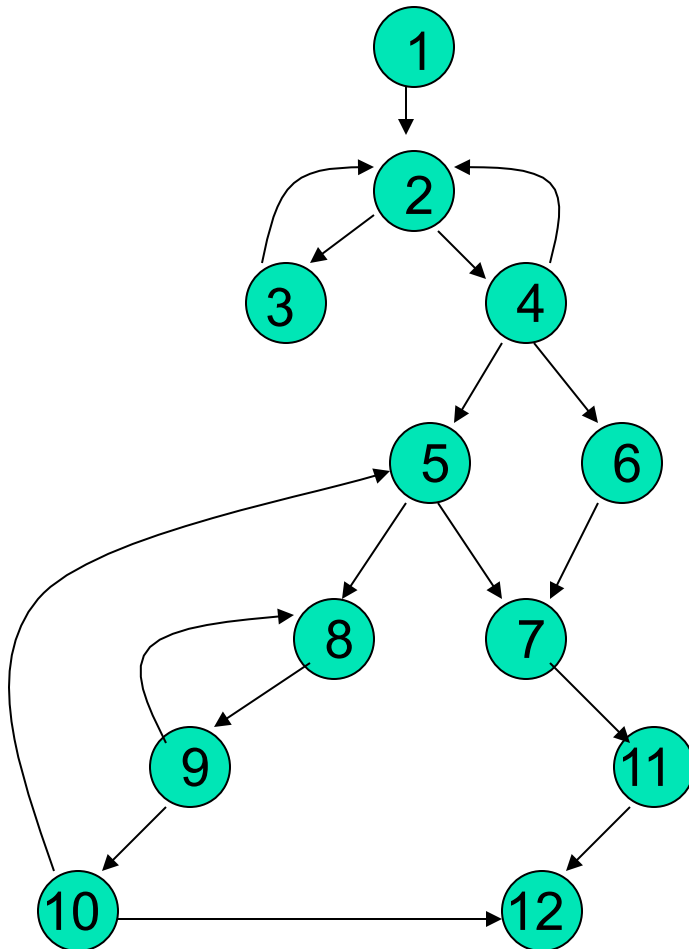
- Say the assignment “ $a \leftarrow b+c$ ” occurs within a loop
 - But “b” and “c” aren't assigned in the loop
 - Would like to move “ $a \leftarrow b+c$ ” somewhere “right in front of” the loop
 - Thereby avoiding redundant reassignments to “a”
- Problem IR/CFG is a directed graph of basic blocks
 - Where, for example, is the “front of” a loop?

A flow graph



Say “ $a \leftarrow b+c$ ” occurs in 9,
To where might we hoist it?

A flow graph



Say “ $a \leftarrow b+c$ ” occurs in 9,
To where might we hoist it?

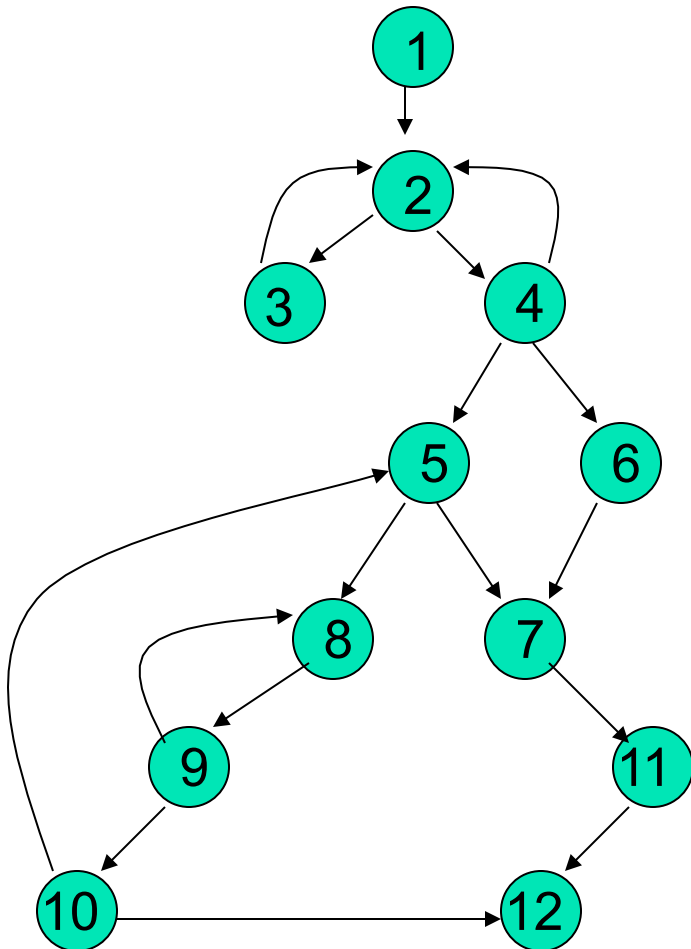
- Nodes 5,4,2,1 seem likely
 - how would I determine that automatically?
 - Impact on other optimizations?
 - Identifying loops with “dominator tree”



Dominators

- Assume we have a flow graph with entry node s_0 with no predecessors
 - I.e., no edge into s_0
- Node d **dominates** n means that
 - d occurs in every path from s_0 to n
- Note that every node dominates itself

A flow graph




- for each pair of nodes d, n
 - does d dominate n ?
 - can think of “ d dominates n ” as a new kind of directed arc in a new graph
- **Question:** is it possible that for $d \neq n$ that:
 - d dominates n , and
 - n dominates d ?

Theorem

Let $d \neq e$ both dominate n . Then, either d dominates e or e dominates d (but not both)

Proof sketch. Assume neither d dominates e nor vice versa.

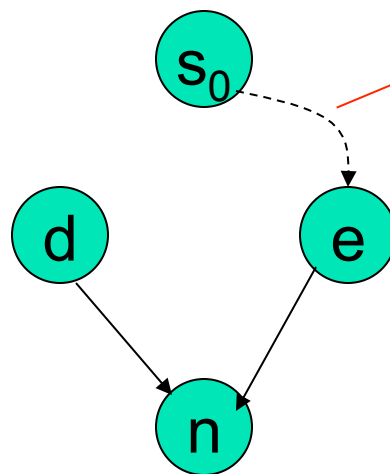


Path to e not containing d exists since d doesn't dominate e

Theorem

Let $d \neq e$ both dominate n . Then, either d dominates e or e dominates d (but not both)



Proof sketch. Assume neither d dominates e nor vice versa.



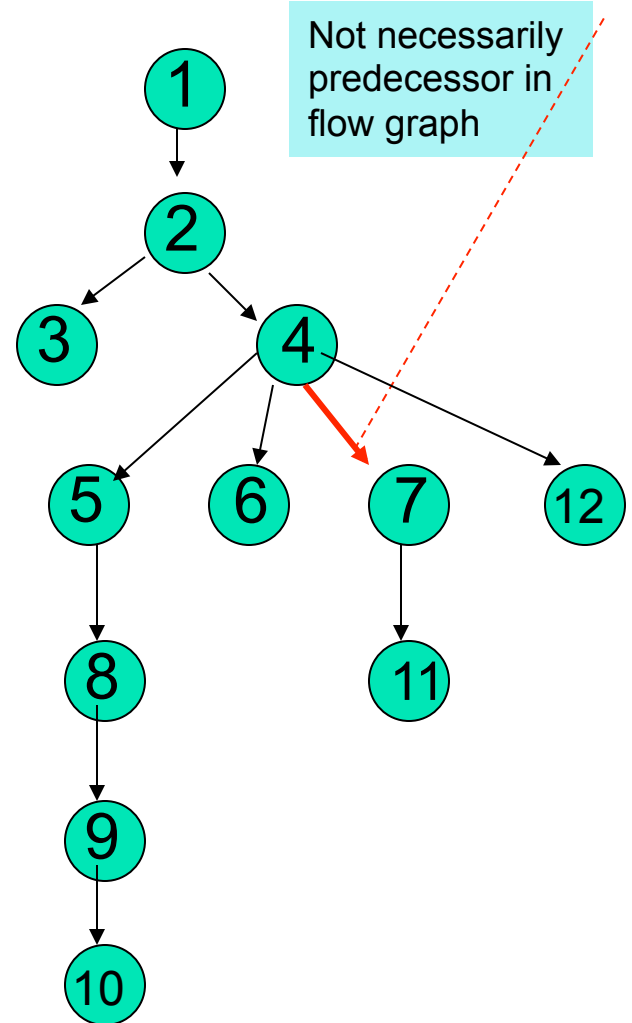
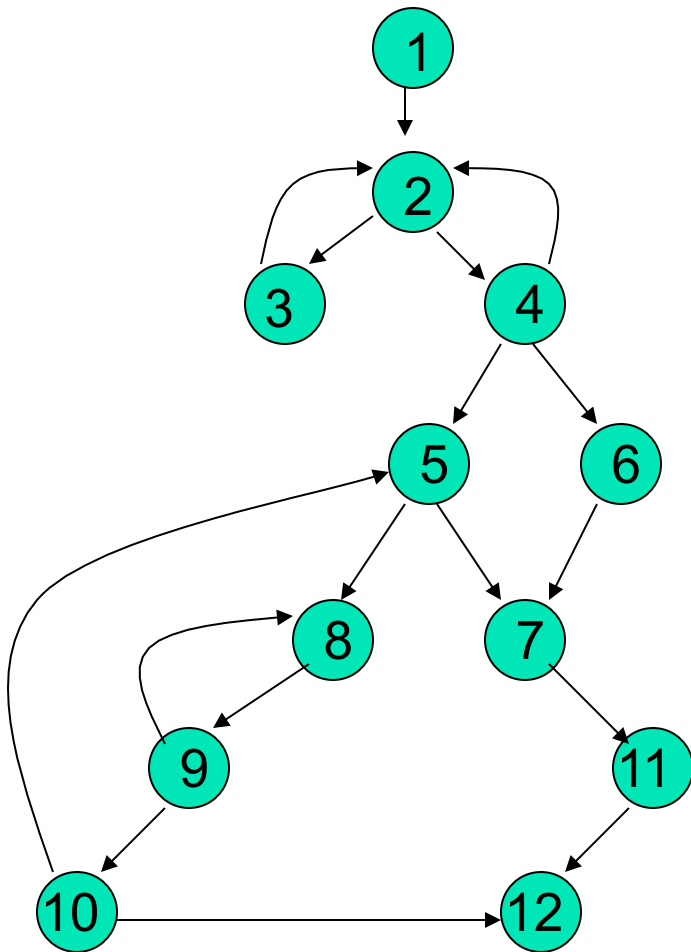
Path to e not containing d exists since d doesn't dominate e

Path from s_0 to n not through d exists. $\therefore d$ doesn't dominate n $\Rightarrow \Leftarrow$

flow graph & its dominator tree



flow graph & its dominator tree





Immediate Dominators

- An **immediate dominator** of node n is a node $\text{idom}(n)$ such that
 - $\text{idom}(n) \neq n$
 - $\text{idom}(n)$ dominates n
 - $\text{idom}(n)$ does not dominate any other dominator of n
- Every node has (at most) one immediate dominator
 - How do we know this?



Calculating Dominators

Let $D[n]$ be the set of nodes that dominate n in a particular flow graph G , we get the following two simultaneous equations*

$$D[s_0] = \{s_0\}$$

$$D[n] = \{n\} \cup \left(\bigcap_{p \in \text{pred}(n)} D[p] \right) \text{ for } n \neq s_0$$

*I wonder how we'd solve equations like this?



Iterative solution

change := true

$D[s_0] := \{s_0\}$

foreach $n \in (\text{Nodes}(G) \setminus \{s_0\})$ { $D[n] := \{n\}$ }

repeat

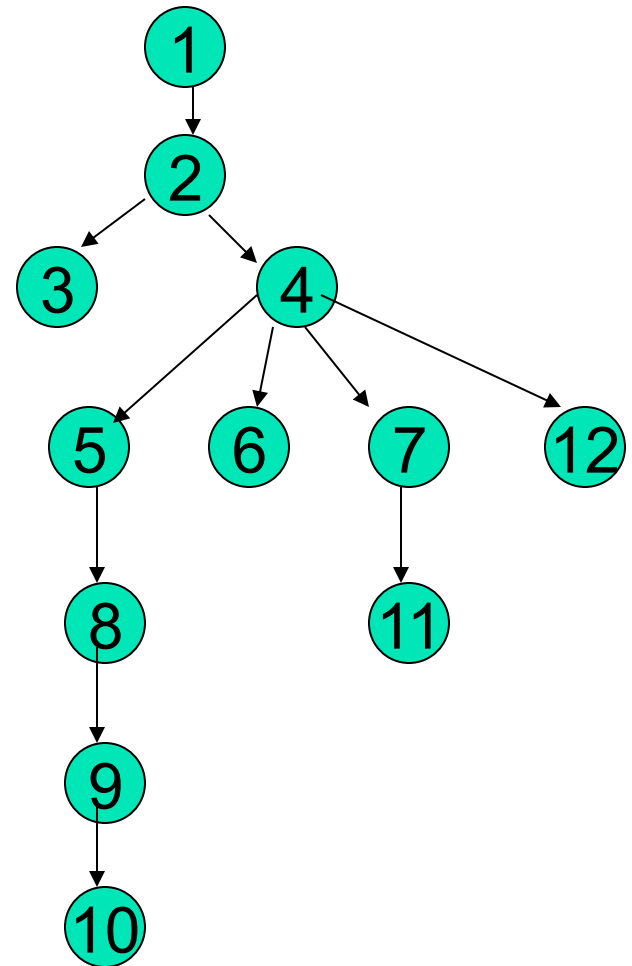
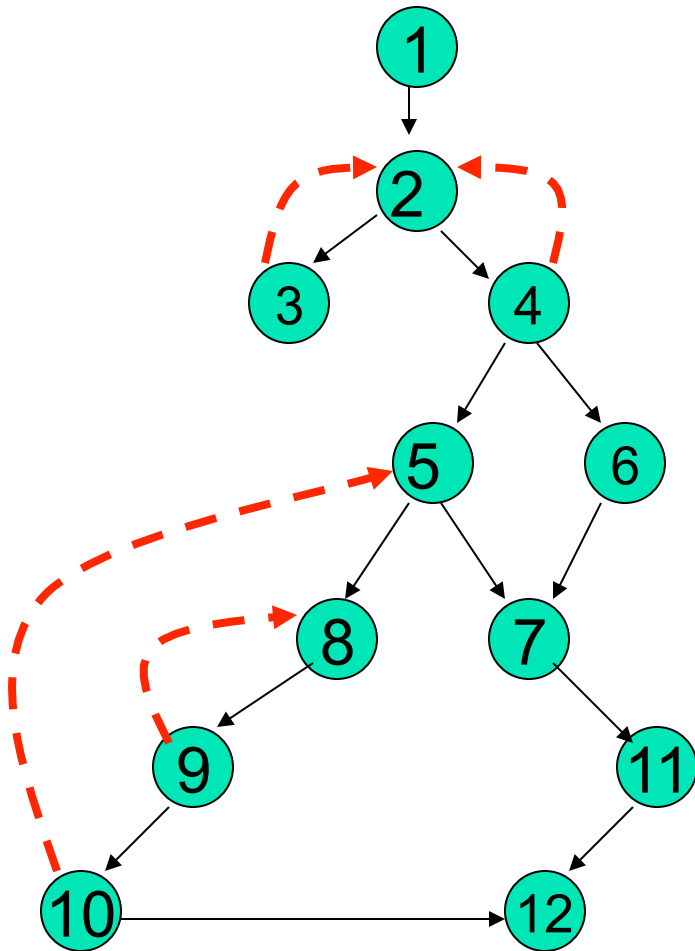
 change := false

 foreach $n \in (\text{Nodes}(G) \setminus \{s_0\})$

$\left(\begin{array}{l} T := \text{Node} \\ \text{foreach } p \in (\text{pred}(n) \setminus \{s_0\}) \{ T := T \cap D[p] \} \\ X := \{n\} \cup T \\ \text{if } X \neq D[n] \text{ then} \\ \quad \text{change} := \text{true} \\ \quad D[n] = X \end{array} \right)$

until (not change)

Backedges are edges $n \rightarrow h$ where h dominates n

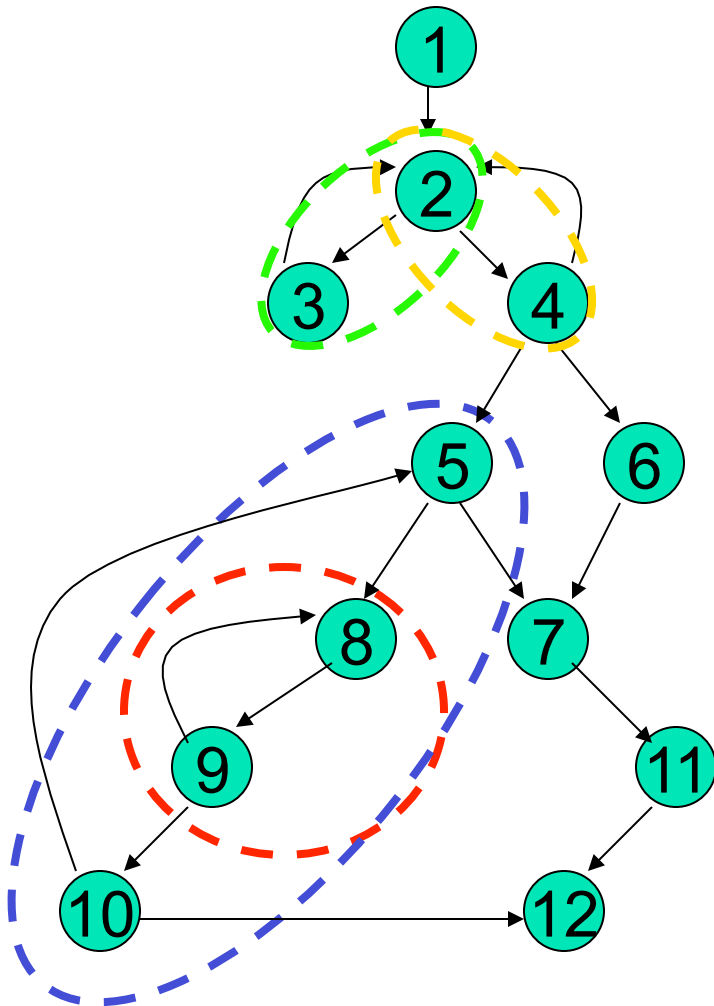




“Natural” Loops

- The natural loop of a backedge $n \rightarrow h$ is
 - The set of x such that
 - h dominates x
 - there is a path from x to n not containing h
- h is called the **header** of this loop

Backedges induce natural loops



The four natural loops are:

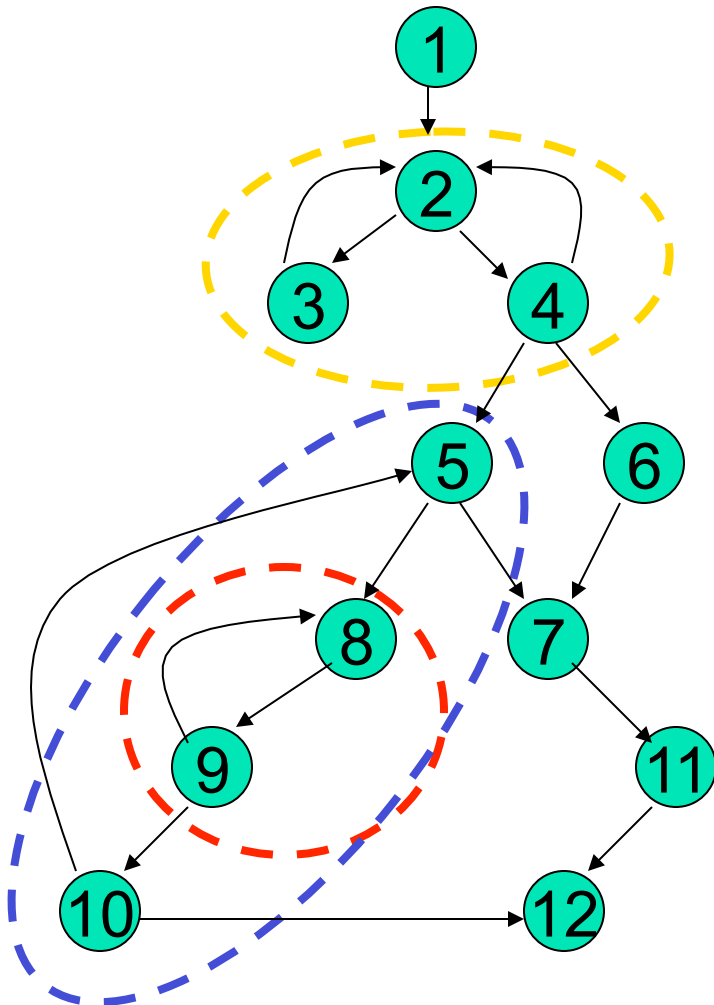
$\{2,3\}$

$\{2,4\}$

$\{5,8,9,10\}$

$\{8,9\}$

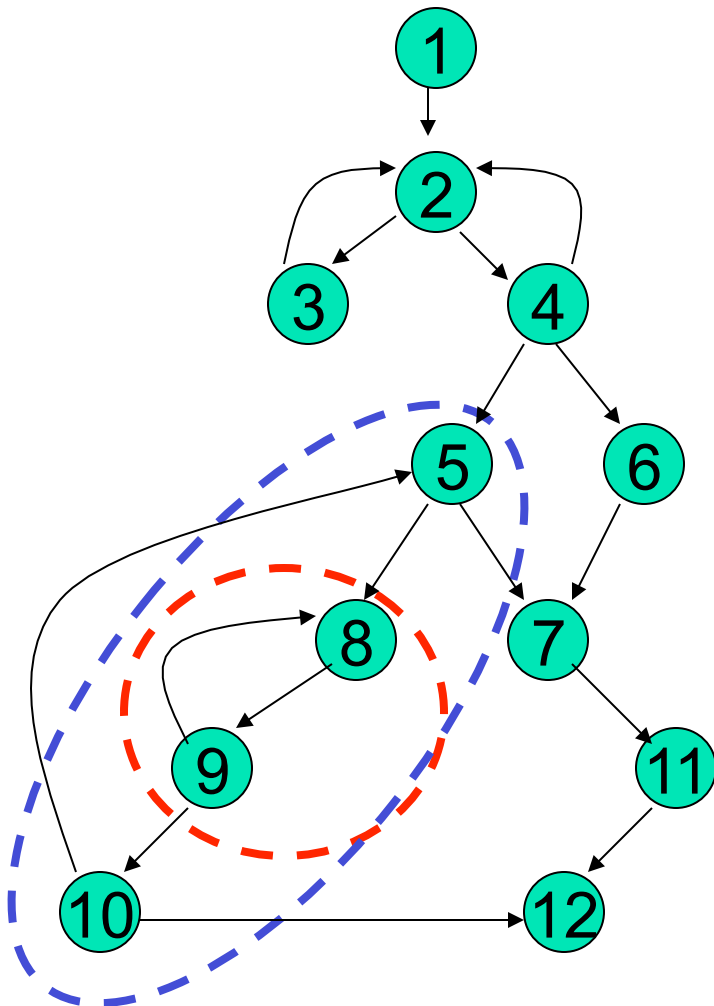
Identifying loops



By merging natural loops with identical headers, one identifies all **loops**

- note that a natural loop is what you think of as a loop
- while a loop inside a compiler is a slightly different animal

Identifying nested loops



Defn: Let A, B be loops with headers a, b s.t. $a \neq b$.
 B is **nested** within A iff $B \subset A$

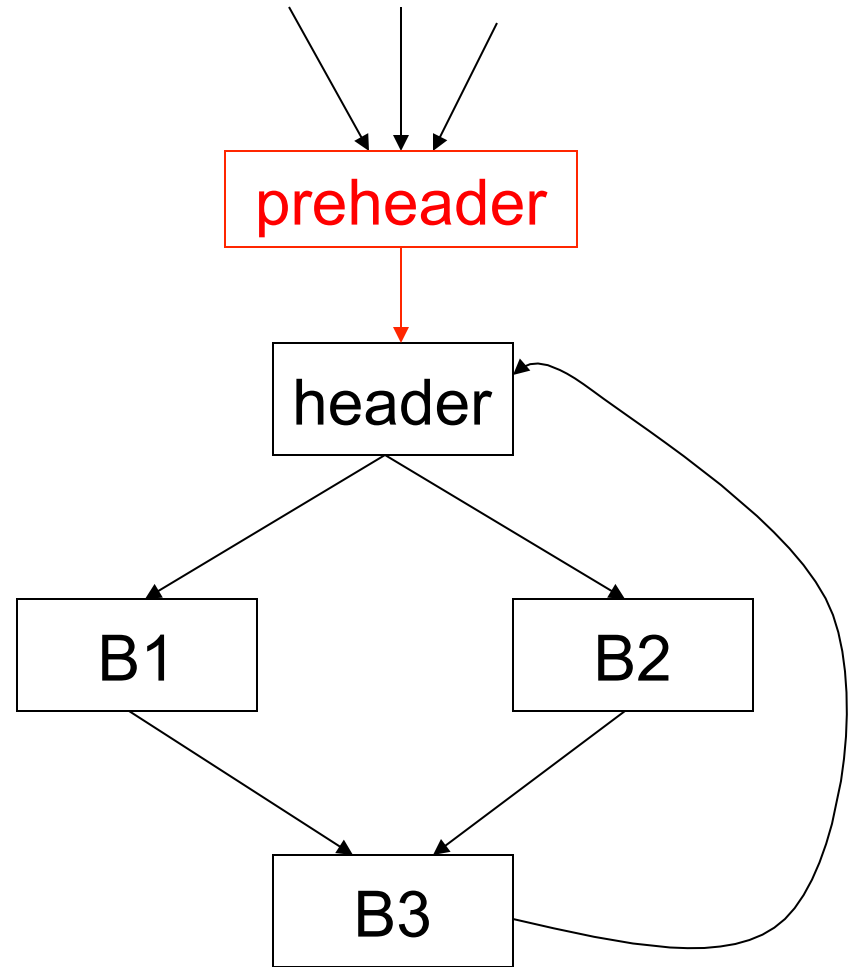
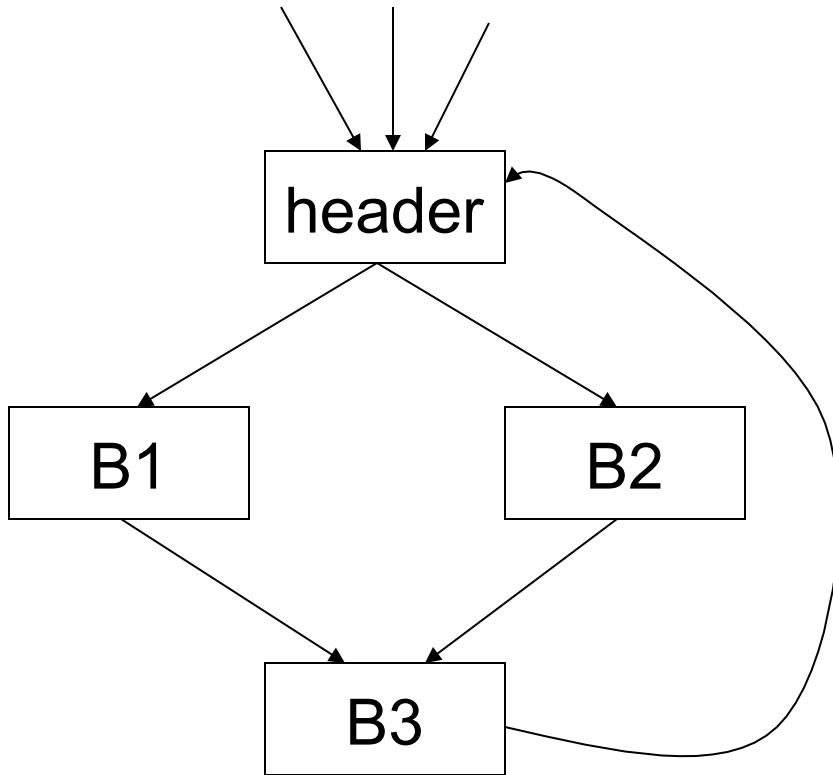
- generally start optimizing inside the innermost loop
- Ex: $\{8, 9\}$ nested within $\{5, 8, 9, 10\}$
- Reminder: “ \subset ” means proper subset
 - i.e., $A \neq B$



Loop pre-header

- Many loop optimizations require moving code from inside a loop to just before its header
- To guarantee that we uniformly have such a place available, we may insert a **loop pre-header**
 - initially empty basic block with a single edge into the header
 - potentially reduces code duplication, among other things
 - the pre-header block dominates the loop (that's important)

Inserting a loop pre-header



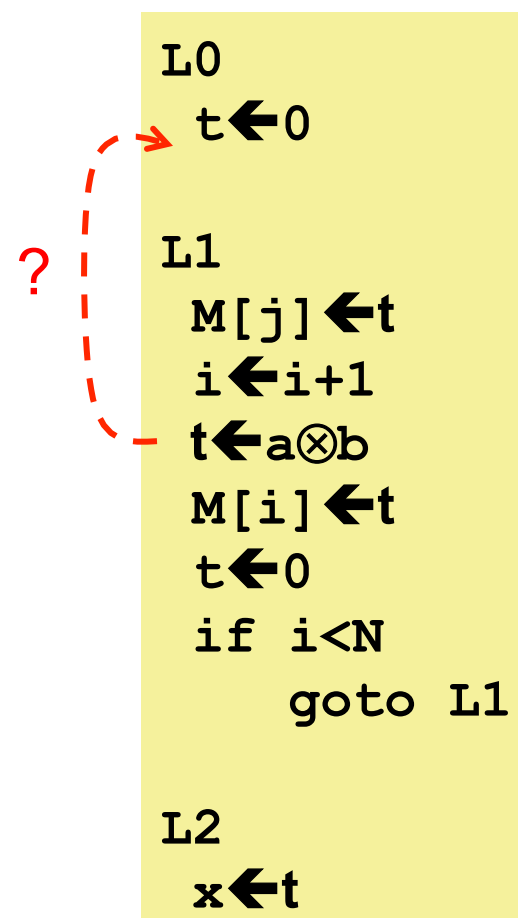
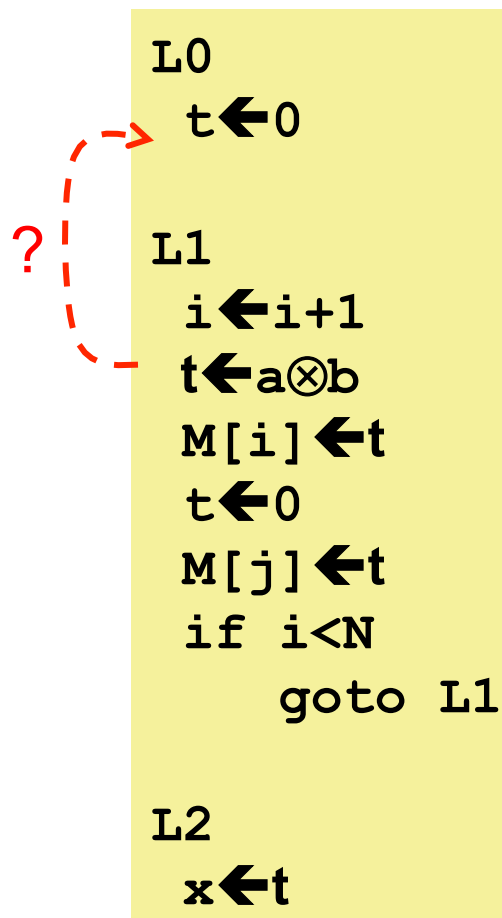
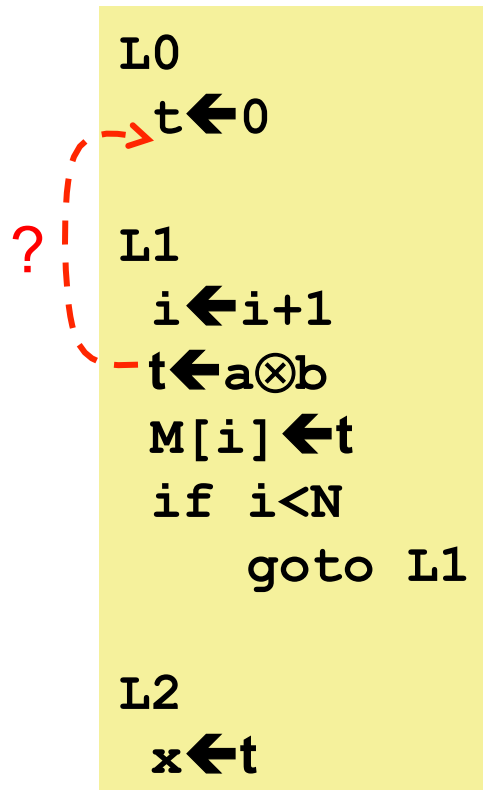


Loop invariant computations

- Let L be a loop and d be “ $t \leftarrow a \otimes b$ ”
 - d is **loop invariant** for L iff
 - “ a ” and “ b ” are constant, or
 - definitions reaching “ a ” and “ b ” occur outside L , or
 - only one definition reaches “ a ” and one reaches “ b ” and they are loop invariant for L
- This is a conservative estimate of loop invariance
 - i.e., it may report d is not loop invariant when it is
 - but it will never say it is loop invariant when it is not

To hoist or not to hoist...

Q: when is it safe to hoist $t \leftarrow a \otimes b$?





Ex 1: To hoist or not to hoist...

Before

```
L0
  t ← 0

L1
  i ← i+1
  t ← a ⊗ b
  M[i] ← t
  if i < N
    goto L1

L2
  x ← t
```

After

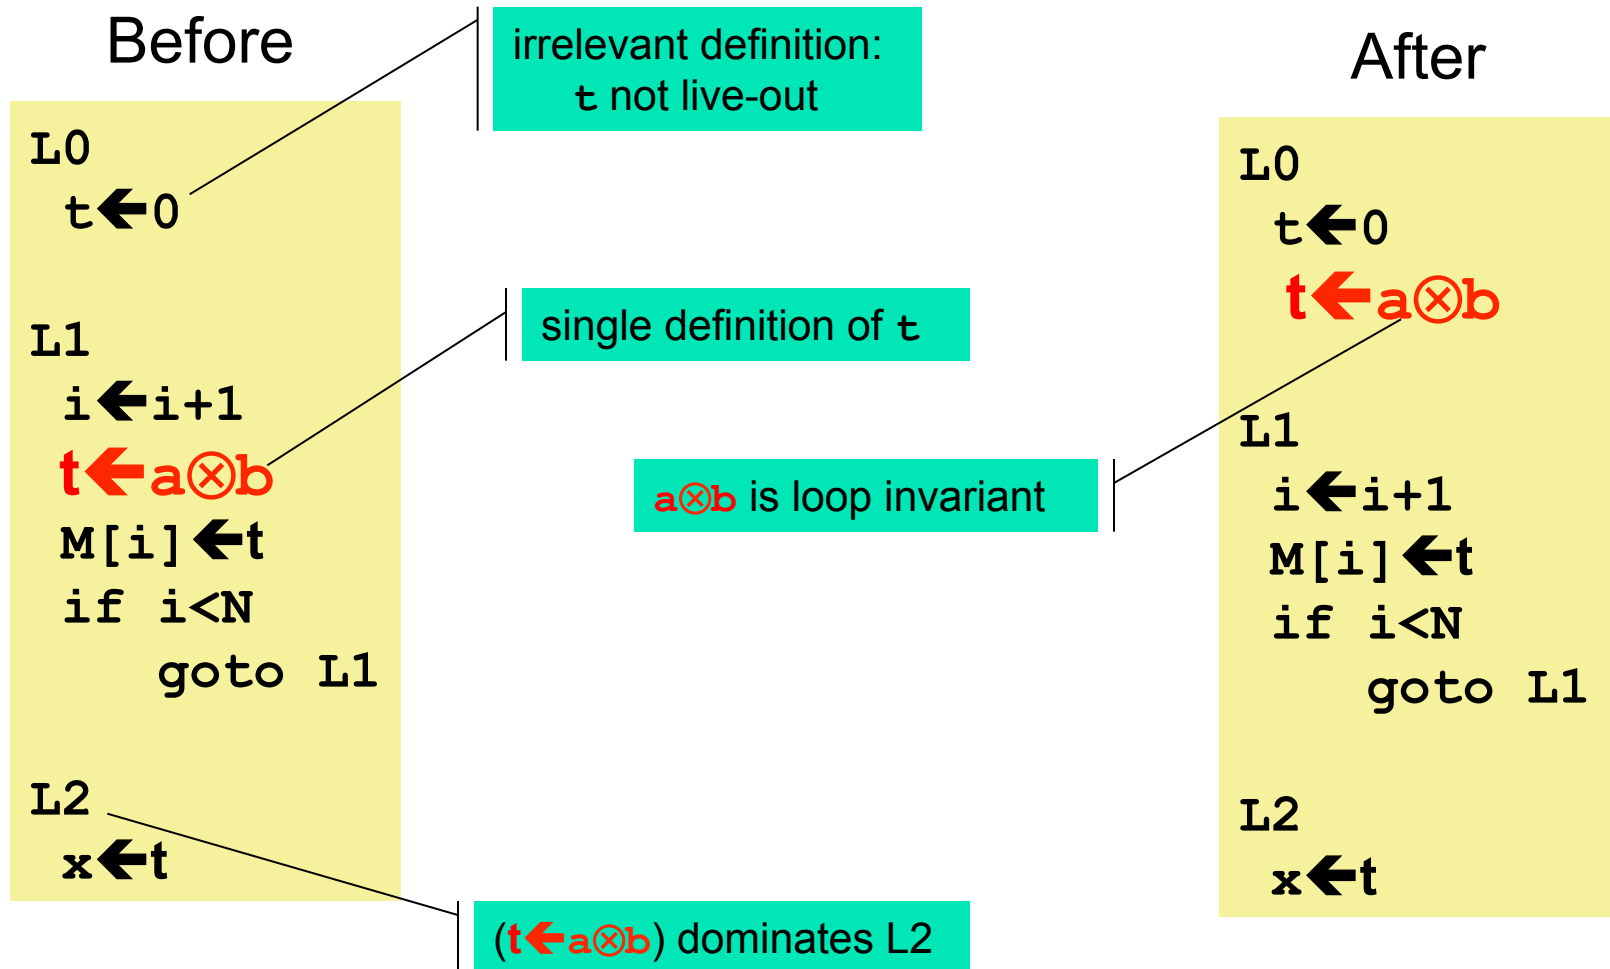
```
L0
  t ← 0
  t ← a ⊗ b

L1
  i ← i+1
  M[i] ← t
  if i < N
    goto L1

L2
  x ← t
```

* We must determine the criteria for hoisting

Determining the criteria for safe hoisting





Ex 2: To hoist or not to hoist...

Before

```
L0
  t ← 0

L1
  branch (i ≥ N) L2
  i ← i + 1
  t ← a ⊗ b
  M[i] ← t
  goto L1

L2
  x ← t
```

After

```
L0
  t ← 0
  t ← a ⊗ b

L1
  branch (i ≥ N) L2
  i ← i + 1
  M[i] ← t
  goto L1

L2
  x ← t
```

To hoist or not to hoist...

Before

```
L0
  t ← 0

L1
  branch (i ≥ N) L2
  i ← i + 1
  t ← a ⊗ b
  M[i] ← t
  goto L1

L2
  x ← t
```

always executes

only executes
when branch fails

can be $a \otimes b$ or 0

After

```
L0
  t ← 0
  t ← a ⊗ b

L1
  branch (i ≥ N) L2
  i ← i + 1
  M[i] ← t
  goto L1

L2
  x ← t
```

always $a \otimes b$



Ex 3: To hoist or not to hoist...

Before

```
L0
  t ← 0

L1
  i ← i+1
  t ← a ⊗ b
  M[i] ← t
  t ← 0
  if i < N
    goto L1

L2
```

After

```
L0
  t ← 0
  t ← a ⊗ b

L1
  i ← i+1
  M[i] ← t
  t ← 0
  if i < N
    goto L1

L2
```

To hoist or not to hoist...

Before

```
L0
  t ← 0

L1
  i ← i+1
  t ← a ⊗ b
  M[i] ← t
  t ← 0
  if i < N
    goto L1
```

always
a ⊗ b

After

```
L0
  t ← 0
  t ← a ⊗ b

L1
  i ← i+1
  M[i] ← t
  t ← 0
  if i < N
    goto L1
```

sometimes
a ⊗ b

Observe: multiple definitions of t inside $L1$ complicate matters

Ex 4: To hoist or not to hoist...

Before

```
L0
  t ← 0

L1
  M[j] ← t
  i ← i+1
  t ← a ⊗ b
  M[i] ← t
  if i < N
    goto L1

L2
  x ← t
```

sometimes $a \otimes b$

After

```
L0
  t ← 0
  t ← a ⊗ b

L1
  M[j] ← t
  i ← i+1
  M[i] ← t
  if i < N
    goto L1

L2
  x ← t
```

always $a \otimes b$



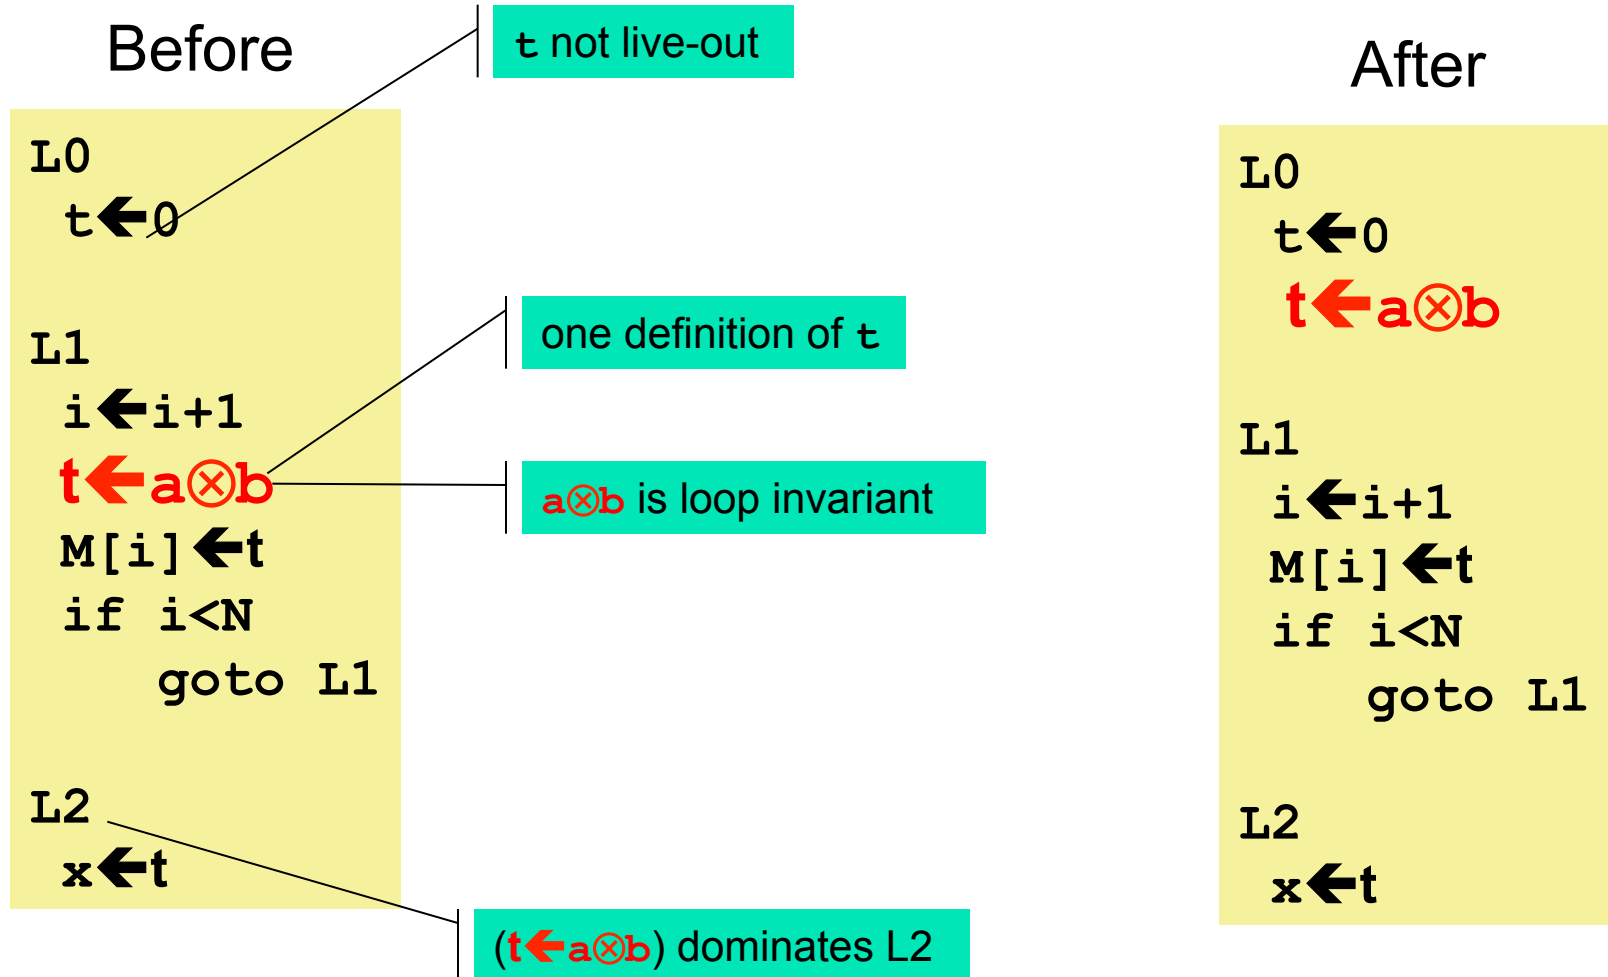
Safe hoisting of $d: t \leftarrow a \otimes b$

- d dominates all loop exits where t is live, and
- only one definition of t in the loop, and
- t is not live-out in the preheader
- assumes d is loop invariant

This is only a conservative approximation!

- will say some hoistings are unsafe when they are safe
- will not say a hoisting is safe when it isn't

To hoist or not to hoist...



*Therefore, hoist!



Induction Variable Analysis

- Some loops have an induction variable
 - a variable “i” that is incremented by a constant or loop invariant amount each iteration
 - for loop-inv. “c”, only definitions of the form:
 - “ $i \leftarrow i + c$ ” or “ $i \leftarrow i - c$ ”
- Other variables may depend entirely on “i”
 - these are called **derived induction variables**
- Identifying induction variables within a loop enables a variety of loop optimizations
 - **strength reduction**: replacing an expensive operation by a less expensive one
 - **induction variable elimination**: removing the variable, thereby (perhaps) shortening the code and reducing register pressure

An Example

```
s ← 0
i ← 0
L1
  branch (i > n) L2
  j ← i * 4
  k ← j + a
  x ← M[k]
  s ← s + x
  i ← i + 1
  goto L1
L2
```

Before

i is an induction variable

- $i * 4$ has values 0, 4, 8, 12, ...
- can perform **strength reduction**

```
s ← 0
i ← 0
j ← 0
L1
  branch (i > n) L2
  j ← j + 4
  k ← j + a
  x ← M[k]
  s ← s + x
  i ← i + 1
  goto L1
L2
```

After



Basic Induction Variables

- Given a loop L with header h
 - “ i ” is a basic induction variable within L iff
 - the only definitions of “ i ” have the form:
 - “ $i \leftarrow i + c$ ” or “ $i \leftarrow i - c$ ”
 - for loop-invariant “ c ”
- Detection of basic induction variables is done by inspecting their form



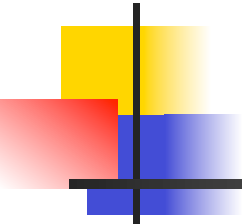
Derived Induction Variables in the family of “i”

- If “i” is an induction variable (basic or otherwise) for loop L, then
 - “j” is a derived induction variable in the family of “i” means
 - all definitions of “j” in L are of the form
 - $j \leftarrow c*i + d$
 - where c,d are loop invariant
 - may be more than one instruction
 - Lingo: j is **determined** by (i,c,d)



Derived induction variables in the family of “i”

- **Key Insight:** definitions of such a “j” may be rewritten w/o reference to “i”
 - That is, replace definition(s) with the effect of “ $j \leftarrow c*i + d$ ” with “ $j \leftarrow j + c*a$ ”
 - where “i” is incremented by “a”
 - c,d --- loop invariant for L
 - N.b., “ $c*a$ ” is either
 - constant: in which case, calculate it
 - loop-inv, but not constant
 - in which case compute in the pre-header



Example: $i \leftarrow i+4, j \leftarrow 2*i+5$

i	$2*i + 5$
0	5
4	13
8	21
12	29
...	...

Initialize	$j \leftarrow 5$
iteration	$j \leftarrow j+2*4$
0 th	5
1 st	13
2 nd	21
3 rd	29
...	...

after each
iteration

$a=4, c=2$



Detecting derived induction variables

Let “j” be an induction variable for L in the family of “i”

- “k” is a derived induction variable for “j” in loop L when:
 - **(Basic)** there is only one definition of “k” in L
 - and that definition is of the form:
 - “k \leftarrow c * j”
 - “k \leftarrow j + d”
 - where c,d are loop invariant



“k” is a derived induction variable for “j” in loop L when:

- **(Basic)** only one definition of “k” in L
 - and that definition is of the form:
 - “k \leftarrow c * j”
 - “k \leftarrow j + d”
 - where c,d are loop invariant
- **(More Complex)**
 - the only definition of “j” that reaches “k” is in the loop
 - i.e., “k” depends entirely on a single definition of “j”
 - and, no definition of “i” occurs on any path between the definitions of “j” and “k”
 - i.e., the dependence of “k” on “i” via “j” maintained



Array bounds checking

“Safe” programming languages
insert dynamic checks to array references
to avoid “out of bounds” references

```
array m[1..100] of int;  
...  
m[i] := 77;  
...
```



```
r1 ← fp + offseti  
branch (r1 > 100) Lerror  
branch (r1 < 1) Lerror  
<code for assignment>  
...
```



Array bounds checking

Under certain circumstances, may be able to determine that the array reference is safe and eliminate the check

- relies on induction variable analysis

```
array m[1..100] of int;  
...  
m[i] := 77;  
...
```



```
r1 ← fp + offseti  
branch (r1 > 100) Lerror  
branch (r1 < 1) Lerror  
<code for assignment>  
...
```




Loop Unrolling Overview

- If the body of a loop L is small, it may be that it spends most of its execution “looping” rather than “computing”; i.e.,
 - incrementing induction variables
 - branching
- Simple example at source level; replace
 - “for ($i=0$; $i++$; $i<2$) { c }” with
 - “ $i=0$; c ; $i=1$; c ”
 - avoids branching, etc.
- The Problem: how do you do this at the machine code/IR level?
- **AKA Software Pipelining**

Unrolling a loop (brute force)

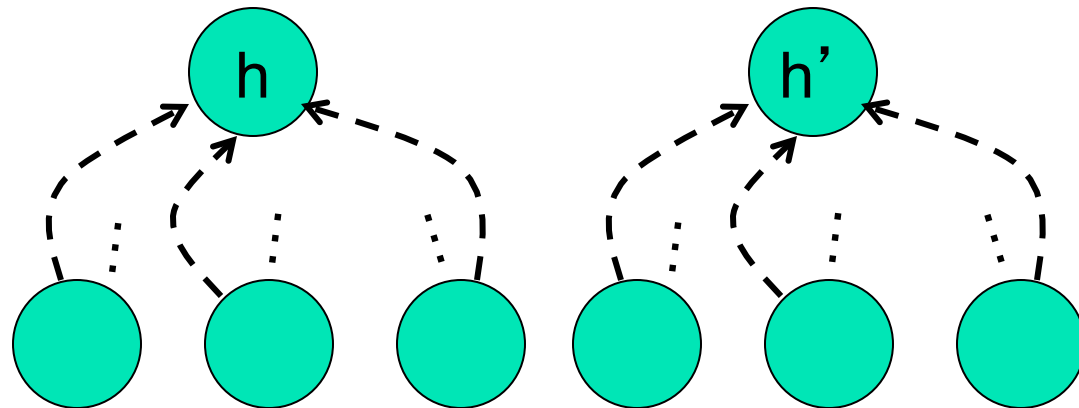
Let L be the loop:



where h is the header and $\cdots\rightarrow$ are backedges

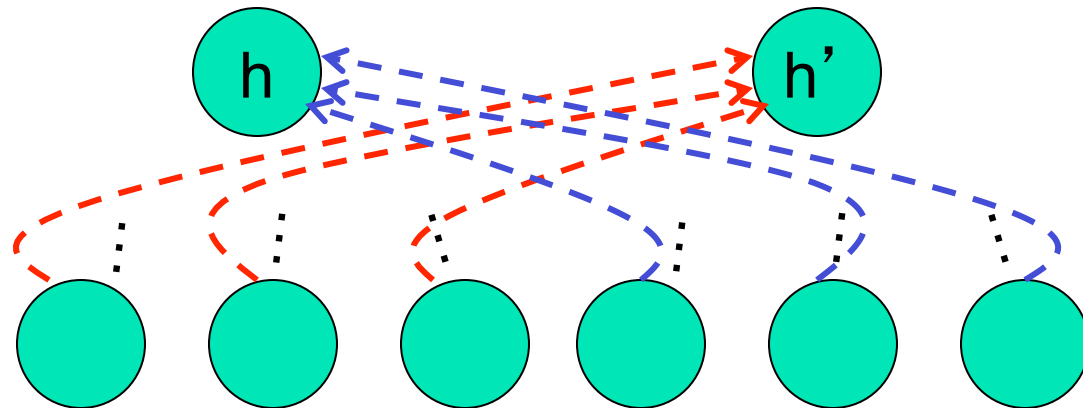
Unrolling a loop (brute force)

Make a copy of L



Unrolling a loop (brute force)

Reroute the backedges





Effects of “brute force” unrolling

Before

```
L1:  x ← M[i]
      s ← s+x
      i ← i+4
      branch (i < n) L1
L2:
```

After

```
L1:  x ← M[i]
      s ← s+x
      i ← i+4
      branch (i ≥ n) L2
L1': x ← M[i]
      s ← s+x
      i ← i+4
      branch (i < n) L1
L2:
```

Q: Does this constitute an improvement?



Empirical Studies show: loops with the following are good candidates for unrolling

- **Single Basic Blocks**

- i.e., straight line code
- with a limited number of floating point & memory operations
 - why? Limits unrolling to loops that are most likely to benefit from instruction scheduling (i.e., ordering code w.r.t. architectural features such as data caches)

- **Small in Length**

- otherwise unrolling may have negative impact on instruction cache performance

- **Simple Loop Control**

- simplifies the unrolling transformation



Typical unrolling candidate

- Note that half of the work is in “looping”
- It's a loop within a single basic block
- few instructions
- has a “repeat-until” structure

```
L1:  x ← M[i]
      s ← s+x
      i ← i+4
      branch (i < n) L1
L2:
```



Fragile* unrolling (K=2)

Before

```
L1:  x ← M[i]
      s ← s+x
      i ← i+4
      branch (i<n) L1
L2:
```

After

```
L1:  x ← M[i]
      s ← s+x
      i ← i+4
      x ← M[i]
      s ← s+x
      i ← i+4
      branch (i<n) L1
L2:
```

**Fragile*: must establish that middle branch may be removed and that loop iterates an even number of times



“Fragile” unrolling (K=2)

Before

```
L1:  x ← M[i]
      s ← s+x
      i ← i+4
      branch (i<n) L1
L2:
```

After

```
L1:  x ← M[i]
      s ← s+x
      i ← i+4
      x ← M[i+4]
      s ← s+x
      i ← i+8
      branch (i<n) L1
L2:
```

...can do better with induction variable analysis



Robust* unrolling (K=2)

Before

```
L1:  x ← M[i]
      s ← s+x
      i ← i+4
      if (i<n) L1 else L2
L2:
```

**Robust*: i.e., works for any number of iterations.

After

```
      if (i<n-8) L1 else L2
L1:  x ← M[i]
      s ← s+x
      x ← M[i+4]
      s ← s+x
      i ← i+8
      if (i<n-8) L1 else L2
L2:  x ← M[i]
      s ← s+x
      i ← i+4
      if (i<n) L2 else L3
L3:
```

Robust* unrolling (K=2)

Before

```
L1:  x ← M[i]
      s ← s+x
      i ← i+4
      if (i < n) L1 else L2
L2:
```

**Robust*: i.e., works for any number of iterations.

After

```
L1:  if (i < n-8) L1 else L2
      x ← M[i]
      s ← s+x
      x ← M[i+4]
      s ← s+x
      i ← i+8
      if (i < n-8) L1 else L2
L2:  x ← M[i]
      s ← s+x
      i ← i+4
      if (i < n) L2 else L3
L3:
```

“twice” block

“once” block

That's all, folks!