

Pattern-driven Reduction in Haskell

William L. Harrison¹

*Pacific Software Research Center
OGI School of Science & Engineering
Oregon Health & Science University
Beaverton, Oregon, USA*

Richard B. Kieburtz²

*Pacific Software Research Center
OGI School of Science & Engineering
Oregon Health & Science University
Beaverton, Oregon, USA*

Abstract

Haskell is a functional programming language with nominally non-strict semantics, implying that evaluation of a Haskell expression proceeds by demand-driven reduction. However, Haskell also provides pattern matching on arguments of functions, in **let** expressions and in the match clauses of **case** expressions. Pattern-matching requires data-driven reduction to the extent necessary to evaluate a pattern match or to bind variables introduced in a pattern. In this paper we provide both an abstract semantics and a logical characterization of pattern-matching in Haskell and the reduction order that it entails.

1 Introduction

Haskell is known as a “lazy” functional language, implying that the default reduction strategy used to evaluate terms to a head-normal form is leftmost-outermost. Furthermore, the data constructors of Haskell’s algebraic data types are non-strict—a constructor application is a head-normal form even if the argument terms are unnormalized.

However, Haskell also employs pattern-matching to select among alternative branches of a **case** expression or alternative equations of a function definition. Pattern-matching can force partial evaluation of an argument or a **case** discriminator, whether or not the value of any variable bound in the

¹ Email: wlh@cse.ogi.edu

² Email: dick@cse.ogi.edu

```

type Name = String
data LS = Lazy | Strict deriving Eq
data P = Pvar Name | Pcondata Name [P] | Ptilde P | Pwildcard           {- Patterns -}
data E = Var Name | Undefined | ConApp (Name,[LS]) [E] | Case E [(P,E)] {- Expressions -}

```

Fig. 1. Abstract Syntax of a Haskell Fragment

pattern is ever demanded. In that sense, the reduction order mandated by pattern-matching is not lazy and may deviate from a leftmost-outermost strategy.

Of course, situations arise in which a programmer “knows” that a certain pattern will occur in an argument, so pattern-matching for control purposes can be shortcut. Haskell provides a prefix notation (\sim) with which a programmer can indicate that pattern-matching is disabled for control. A control-disabled pattern (called *irrefutable* in the Haskell literature) serves only to specify that variables occurring in the pattern are bound by a successful match. These bindings can be evaluated lazily. Control-disabled patterns call for a reduction strategy different from nominal, pattern-driven reduction.

The reduction strategy employed to evaluate Haskell programs containing pattern constructs is implicit in the language definition, although a careful reading is necessary. It is perhaps the language aspect least well understood by Haskell amateurs. This paper gives a semantics for the fragment of Haskell that involves pattern-matching. It also proposes inference rules of a programming logic for this language fragment. The semantics and the logic provide complementary but coherent formal descriptions that are consistent with the language definition [4]. They should help a reader to comprehend the nuances of Haskell reduction strategies.

2 A Haskell fragment and its informal semantics

2.1 Patterned abstractions

Patterns may occur in several different syntactic contexts in Haskell—in **case** branches, explicit abstractions, or on the left-hand sides of definitions. We say that a pattern is *abstracted* if it occurs in an operand position on the left-hand side of a function definition³, under a lambda-symbol (the backslash, in Haskell) or to the left of the arrow symbol (\rightarrow) in a case branch. Since the roles played by abstracted patterns are similar in every context, we shall focus on patterns as they occur in **case** expressions.

2.1.1 Evaluating case expressions

When a **case** expression is evaluated, the first case branch is applied to the case discriminator (the expression between the keywords **case**...**of**). If the

³ In a local (**let**) definition, a pattern may occur as the entire left-hand side of an equation. Such an occurrence is implicitly control-disabled, even if it is not prefixed by the character (\sim).

case discriminator matches the abstracted pattern of this branch, then the body of the case branch is evaluated in a context extended with the value bindings of pattern variables made by the match. If the discriminator fails to match the pattern, then the next in the list of case branches is applied to the discriminator, and so on, until a pattern is found to match. If there is no branch whose pattern matches, then evaluation of the case expression fails with an unrecoverable error.

2.1.2 Matching abstracted patterns

An abstracted pattern fulfills two roles:

- **Control:** A **case** discriminator expression is evaluated to the extent necessary to determine whether it matches the pattern of a case branch. If the match fails, control shifts to try a match with the next alternative branch, if one is available.
- **Binding:** When a match succeeds, each variable occurring in the pattern is bound to a subterm corresponding in position in the (partially evaluated) **case** discriminator. Since patterns in Haskell cannot contain repeated occurrences of a variable, the bindings are unique at any successful match.

2.1.3 Variables and wildcard patterns

A variable is itself a pattern which matches any term⁴. Thus a match with a variable never fails and always accomplishes a binding. A term need not be evaluated to match with a pattern variable.

Haskell designates a so-called wildcard pattern by the underscore character (`_`). The wildcard pattern, like a variable, never fails to match but it entails no binding.

2.1.4 Constructor patterns: *strict* and *lazy*

When a data constructor occurs in a pattern, it must appear in a *saturated* application to sub-patterns. That is, a constructor typed as a k -ary function in a datatype declaration must be applied to exactly k sub-patterns when it is used in a pattern.

When a constructor occurs as the top-level operator in a pattern, a match can occur only if the **case** discriminator evaluates to a term that has the same constructor as its primary operator. Subterms of the discriminator must match the corresponding sub-patterns of the constructor pattern or else the entire match fails. If a sub-pattern happens to be a variable or a wildcard, no further evaluation of the corresponding sub-term of the matching expression is required.

However, a constructor may be declared (in a datatype declaration) to be *strict* in one or more of its argument positions by prefixing the character (`!`) to the type expressions in these argument positions. When a constructor is strict

⁴ As Haskell is strongly typed, a variable can only be compared with terms of the same type.

```

-- Semantic Functions for E and P                                -- Environments
mE :: E -> Env -> V                                           type Name = String
mP :: P -> V -> Maybe [V]                                       type Env = Name -> V

-- Domain of Values      {- functions -}      {- structured data -}
data V =      FV (V -> V)      |      Tagged Name [V]

-- Function composition (diagrammatic)                -- Kleisli composition (diagrammatic)
(>>>) :: (a -> b) -> (b -> c) -> a -> c                (<>) :: (a->Maybe b)->(b->Maybe c)-> a->Maybe c
f >>> g = g . f                                           f <> g = \ x -> f x >>= g

-- Domains are pointed                                    -- Purification: the "run" of Maybe monad
bottom :: a                                               purify :: Maybe a -> a
bottom = undefined                                       purify (Just x) = x
                                                         purify Nothing  = bottom

-- Alternation                                           -- Semantic "seq"
fatbar :: (a->Maybe b) ->                                semseq :: V -> V -> V
      (a->Maybe b) -> (a->Maybe b)                       semseq x y = case x of
f 'fatbar' g = \ x -> (f x) 'fb' (g x)                    (FV _)      -> y ;
  where fb :: Maybe a -> Maybe a -> Maybe a                (Tagged _ _) -> y
      Nothing 'fb' y = y
      (Just v) 'fb' y = (Just v)

```

Fig. 2. Semantic Operators

in its i^{th} argument position, an application of the constructor will evaluate its i^{th} argument to head normal form. Thus a pattern match involving a constructor declared to be strict in one or more argument positions implicitly forces evaluation of the corresponding subexpressions of the matching term.

2.1.5 Control-disabled patterns

Disabling a pattern for control does not disable the binding function of a match, it merely defers binding until further computation demands a value for one of the variables occurring in the pattern. When that happens, the focus of computation returns to the deferred pattern match, which is fully computed in order to bind the variables introduced in the pattern. Should a deferred pattern match fail, no alternative is tried, as might have been the case in a normal match failure. Failure of a deferred pattern match causes an unrecoverable program error.

3 Formal Semantics

This section outlines the formal semantics of the Haskell fragment considered in this paper. This semantics has been described in detail elsewhere [3], so the presentation here will be brief. The semantics is presented as a metacircular interpreter for the Haskell fragment whose abstract syntax is given in Figure 1. The interpreter, written in Haskell itself, makes use of standard techniques and structures from the denotational description of programming languages and uses monads to model pattern-matching.

Although the semantic metalanguage here *is* Haskell, care has been taken to use notation which will be recognizable by any functional programmer. However unlike many functional languages, Haskell has explicit monads, and

```

mP :: P -> V -> Maybe [V]
mP (Pvar x) v = Just [v]
mP (Pconddata n ps) (Tagged t vs) = if n==t then (stuple (map mP ps) vs) else Nothing
mP Pwildcard v = Just []
mP (Ptilde p) v = Just(case mP p v of { Nothing -> replicate lp bottom ; Just z -> z })
    where lp = length (fringe p)
          replicate 0 x = []
          replicate n x = x : (replicate (n-1) x)

fringe :: P -> [Name]
fringe (Pvar n) = [n]
fringe (Ptilde p) = fringe p
fringe Pwildcard = []
fringe (Pconddata _ ps) = concat (map fringe ps)

stuple :: [V -> Maybe [V]] -> [V] -> Maybe [V]
stuple [] [] = Just []
stuple (q:qs) (v:vs) = do { v' <- q v ; vs' <- stuple qs vs ; Just (v'++vs') }

```

Fig. 3. Semantics of a Haskell Fragment: Patterns

so we give an overview here of Haskell’s monad syntax⁵. The semantics makes use of an error monad [6], which is modeled in Haskell by the `Maybe` monad. This monad is based upon the datatype `Maybe a` as given below, in which the constructor `Just` encloses a normal expression of type `a` and the constructor `Nothing` models an error condition as a data value.

In every monad there is a unit operator (called `return` in Haskell) and a binary operation which is called “bind” in Haskell and represented by the infix operator (`>>=`). The unit of a monad injects an ordinary (non-monadic) value into the structure of the monad. The bind operation is like a function application expressed in diagrammatic order (i.e. *arg* `>>=` *fn*). It accounts for propagation of the monadic structure of the argument through the computation of the application, which may affect that structure. The unit and bind operators satisfy a set of equations that hold for all monads, thus providing a uniform algebraic framework in which to express computational effects.

The structure of Haskell’s `Maybe` monad is specified by:

```

data Maybe a = Just a | Nothing
return :: a -> Maybe a
return = Just
(>>=) :: Maybe a -> (a -> Maybe b) -> Maybe b
(Nothing >>= f) = Nothing
(Just x >>= f) = f x

```

Haskell provides an alternative syntax for bind (`>>=`) called “do notation” which is defined by:

$$\text{do } \{ y \leftarrow x ; f \} = (x \gg= (\backslash y \rightarrow f))$$

Figure 2 contains a description of the semantic setting for the Haskell fragment considered in this paper. It is in most respects a conventional denotational-style semantics for a functional language. The dynamic semantics for expressions, `mE`, maps an expression and an environment to a value in the domain of values `V`. This domain `V` is itself structured conventionally as a universal domain construction [2] over the sum of functions and structured data. The semantics of patterns, `mP`, takes a pattern and returns a map of type `(V->Maybe [V])` (more will be said about this type in Section 3.1). We use

⁵ We assume the reader has some familiarity with monads [6].

two composition operators, both written as infix operators in diagrammatic order. The symbol (\ggg) denotes function composition and the symbol ($\langle \rangle$) denotes Kleisli composition in the *Maybe* monad. It is assumed that the domain is *pointed* in every type; that is, every domain contains a bottom element **bottom** (usually written \perp), which is modeled in Haskell by the polymorphic constant **undefined**.

Figure 2 also displays two combinators integral to modeling **case** expressions and patterns, called **fatbar** and **purify**. If **m1** and **m2** have type $(V \rightarrow \text{Maybe } V)$, then $((\text{fatbar } m1 \ m2) \ v)$ exhibits sequencing behavior similar to $(\text{case } v \text{ of } \{ m1 ; m2 \})$. The **purify** operator converts a *Maybe*-computation into a value, sending a **Nothing** to **bottom**. Post-composing with **purify** signifies that expressions whose evaluation produces certain pattern-match failures (e.g., exhaustion of the branches of a **case** expression) ultimately denote **bottom**.

3.1 Dynamic Semantics of Haskell Pattern Fragment

In the evaluation of $(\text{case } e \{ p_1 \rightarrow e_1; \dots; p_n \rightarrow e_n \})$, patterns p_i are matched against e in left-to-right order until a successful match p_s is found. Then, the value of the whole expression is the value of the expression e_s . If no such match is found, then the value of the whole expression is undefined. For example, evaluating the following **case** expressions gives these results:

```
data Tree = T Tree Tree | S Tree | L | R
case T L R of {T (S x) y -> y; T x y -> x}      ----> L
case T L R of {T ~(S x) y -> y; T x y -> x}      ----> R
case T L R of {T ~(S x) y -> x; T x y -> y}      ----> program error (match)
case T L R of {~(T (S x) y) -> y; T x y -> x}    ----> program error (match)
```

In the first of the **case** expressions above, the constructor **L** fails to match the embedded pattern $(S \ x)$ in the first case branch. The match failure shifts control to the second case branch. In the second example, the embedded pattern $\sim(S \ x)$ is control-disabled. The term $(T \ L \ R)$ thus matches the pattern $(T \ \sim(S \ x) \ y)$ binding **R** to the variable **y**. In the third example, the body of the first case branch demands a value for **x**, thereby forcing a deferred match of the subterm **L** with the pattern $\sim(S \ x)$. The deferred match fails, resulting in a program error. The fourth example illustrates that a deferred match of the term $(T \ L \ R)$ against the pattern $(T \ (S \ x) \ y)$ fails, although the match was evaluated in response to a request for a binding for **y** alone. Generally, the effect of deferring pattern match failure is characterized by the following equivalence: $(mP \ (\sim p) \ v)$ is **Just**[**bottom**, ..., **bottom**] $\Leftrightarrow (mP \ p \ v)$ is **Nothing**. Even when $(mP \ p \ v)$ fails (i.e., is **Nothing**), $(mP \ (\sim p) \ v)$ still succeeds, but all of the bindings created thereby are **bottom**.

4 Logic

While the dynamic semantics defines a meaning for expressions by providing an abstract evaluation model, a verification logic expresses static assertions

```

mE :: E -> Env -> V
mE (Var n) rho      = rho n
mE (Case e ml) rho  = mcase rho ml (mE e rho)
mE (ConApp (n,ls) es) rho = evalL (zip es ls) rho n []
  where evalL :: [(E,LS)] -> Env -> Name -> [V] -> V
        evalL [] rho n vs      = Tagged n vs
        evalL ((e,Strict):es) rho n vs = semseq (mE e rho)
                                          (evalL es rho n (vs ++ [mE e rho]))
        evalL ((e,Lazy):es) rho n vs  = evalL es rho n (vs ++ [mE e rho])
mE Undefined rho    = bottom

match :: Env -> (P,E) -> V -> Maybe V
match rho (p,e) = mP p <> ((\vs -> mE e (extend rho xs vs)) >>> Just)
  where xs = fringe p

mcase :: Env -> [(P,E)] -> V -> V
mcase rho ml = (fatbarL $ map (match rho) ml) >>> purify
  where fatbarL :: [V -> Maybe V] -> V -> Maybe V
        fatbarL ms = foldr fatbar (\ _ -> (Just bottom)) ms

```

Fig. 4. Semantics of a Haskell Fragment: Expressions

about properties of the semantics. An assertion can take the form of a k -ary predicate applied to k terms. For simplicity, we restrict ourselves here to unary predicates ($k = 1$).

We write $t :: P$ for the assertion that term t satisfies predicate P . Because function and data constructor applications are non-strict in Haskell's evaluation semantics, two notions of satisfaction of a predicate are sensible.

We say that a predicate, P , is *weakly* satisfied by an expression t if the denotation of t belongs to the set specified by P . It is *strongly* satisfied if, in addition, the denotation of t is not the bottom element in its type domain. By convention, a predicate assumes its weak interpretation unless otherwise annotated. An otherwise weak predicate may be explicitly strengthened by prefixing the symbol (\$).

P -logic⁶, a verification logic for Haskell, is based upon the familiar Gentzen-style sequent calculus [1]. In this section we give a brief introduction to the syntax of P -logic, as well as some inference rules that are relevant to pattern-matching in Haskell.

4.1 Predicates in P -logic

Atomic, unary predicates include the predicate constants, **Univ** and **UnDef**, which are respectively satisfied by all terms and by only those terms whose values are undefined.

There are two principal ways that compound predicates are formed in P -logic.

- (i) The constructors of datatypes declared in a Haskell program are implicitly “lifted” to act as predicate constructors in P -logic. For example, in the context of a program, the list constructor $(:)$ combines an expres-

⁶ The name P -logic is short for *Programatica logic*, as the logic has been developed as part of the Programatica project [5].

sion h of type a and an expression t of type $[a]$ into a new expression $(h : t)$ of type $[a]$. In the context of a formula, the same constructor combines a predicate P and a predicate Q into a new predicate, $(P : Q)$. This predicate is satisfied by a Haskell expression that normalizes to a term of the form $(h : t)$ and whose component expressions weakly satisfy the assertions $h :: P$ and $t :: Q$. The default mode of interpretation of the component predicates is weak because the semantics of the data constructor does not require evaluation of its arguments.

- (ii) The “arrow” predicate constructor is used to compose predicates that express properties of case branches. An arrow predicate $P \rightarrow Q$ is satisfied by a case branch $(p \rightarrow e)$ if, whenever the case discriminator satisfies the pattern predicate, P , the body, e , of the case branch satisfies Q .

4.2 Inference Rules for Properties of the Haskell Fragment

4.2.1 Constructor application

Rules for constructor application are derived from a Haskell datatype declaration. A datatype declaration serves to define the data constructors of the type, giving the signature of each constructor as a sequence of type expressions.

$$\text{data } T = \dots \mid C^{(k)} \tau_1 \dots \tau_k \mid \dots$$

Notice that the strictness annotations from the signature of a constructor are represented explicitly in the abstract syntax of a constructor application, although they are not manifested in the concrete syntax.

A constructor application is lifted to a predicate constructor application by the function:

```
sigma :: E -> [Pr] -> Pr
sigma (ConApp (n,ls) _) prs =
  let prs' = take (length ls) prs
      s = and $ map (\(pr,l) -> isStrong pr || l==Lazy) (zip prs' ls)
      where isStrong (Strong _) = True
            isStrong _ = False
  in if s then Strong $ ConPred (n,ls) prs'
      else ConPred (n,ls) prs'
```

where ls lists the strictness declaration (**Lazy** or **Strict**) of the constructor in each argument position. Notice that when a predicate constructor lifted from a strict data constructor is applied to a predicate argument, the resulting predicate is strong only if the argument predicate is so, whereas a predicate derived from a lazy data constructor is always strong. A strong predicate formula $\$C^{(k)} P_1 \dots P_k$ is satisfied by a well-defined term of the form $C^{(k)} t_1 \dots t_k$ whenever each of the t_j satisfies the corresponding predicate P_j .

The rule schemas for properties of a constructor application are:

$$(1) \quad \frac{\Gamma \vdash_{\mathcal{P}} t_1 :: P_1 \dots \Gamma \vdash_{\mathcal{P}} t_k :: P_k}{\Gamma \vdash_{\mathcal{P}} C^{(k)} t_1 \dots t_k :: C^{(k)} P_1 \dots P_k} \quad (1 \leq k)$$

```

-- Abstract syntax of predicates
data Pr = Univ | UnDef | PrCon Name [Pr] | Strong Pr

-- Skeleton of a pattern
skeleton :: P -> (Name -> Pr) -> Bool -> (Pr, Bool)
skeleton (Pvar n) e b = case (e n) of
    pr@(Strong _) -> (pr, True)
    pr             -> (pr, False)
skeleton Pwildcard e _ = (Univ, False)
skeleton (Pconddata (n, ls) ps) e b =
    let (prs, s) = mapSkel ps e True
    in if b || s then (Strong (PrCon n prs), True)
    else (Univ, False)
skeleton (Ptilda p) e _ = skeleton p e False

mapSkel :: [P] -> (Name -> Pr)
        -> Bool -> ([Pr], Bool)
mapSkel [] _ _ = ([], False)
mapSkel (p:ps) e b =
    let (pr, s) = skeleton p e b
    (prs, s') = mapSkel ps e b
    in ((pr:prs), s || s')

```

Fig. 5. Calculation of Pattern Predicates

and, where C and K are distinct constructors in the same data type:

$$(2) \quad \frac{}{\Gamma \vdash_{\mathcal{P}} C^{(k)} t_1 \dots t_k :: \$\neg K^{(n)} \underbrace{\text{Univ} \dots \text{Univ}}_{n\text{-times}}}$$

4.3 Pattern matching

Match clauses have associated with them predicates of a distinct kind. A match clause whose expression body has the Haskell type (τ) may satisfy a predicate of type *Maybe_Pred* τ . These predicates are formed either with the unary predicate constructor *Just* or the nullary constructor *Nothing*.

4.3.1 Pattern predicates

Because patterns may be nested to arbitrary depths, it is inconvenient to use the syntax of patterns directly in rules. Instead we define a syntactically flattened representation for patterns to allow a simpler representation of pattern predicates in rules.

First, we need the concept of the *fringe* of a pattern, the variables that have defining occurrences in the pattern, listed from left to right.

Definition 4.3.1.1: The *fringe* of a pattern is the list of (distinct) variables whose definition is given by induction on the abstract syntax of patterns by the Haskell function `fringe` in Figure 5.

Next, we define a function that produces a pattern predicate from a pattern and an environment that binds predicates to the variables in the fringe of the pattern. Note that

$$\text{fst}(x, y) = x \text{ and } \text{zip} [a_1, \dots, a_n] [b_1, \dots, b_n] = [(a_1, b_1), \dots, (a_n, b_n)]$$

Definition 4.3.1.2: The *pattern predicate* formed by instantiating a pattern relative to a predicate environment is defined inductively by the Haskell function `skeleton` in Figure 5.

We use the notation $\pi(p)$ to designate a “flattened” pattern predicate con-

structor. Formally,

$$\pi(p) P_1 \cdots P_n =_{\text{def}} \text{fst } (\text{skeleton } p \text{ (zip (fringe } p) [P_1, \dots, P_n]) \text{ True})$$

where $n = \text{length } (\text{fringe } p)$

□

The pattern predicate calculated by *skeleton* will ignore control-disabled subpatterns that occur in a pattern, replacing them by the universal predicate, *Univ*, *unless* an explicitly strengthened predicate is bound in the environment to a variable in the fringe of such a subpattern. In such a case, the skeleton of the subpattern is fully elaborated in the *skeleton* computation. In consequence, if an instance of a verification rule such as Rule (3) uses strong predicates in its hypotheses, then the pattern predicate in its conclusion will require a pattern match that evaluates all subterms that are asserted by the strong predicates to be well-defined.

For example, two pattern predicates that are derived from one of the patterns given in the example of Section 3.1 are:

$$\begin{aligned} \pi(\text{T } \sim(\text{S } x) \text{ y}) \text{ Univ Univ} &= \$(\text{T Univ Univ}) \\ \pi(\text{T } \sim(\text{S } x) \text{ y}) \$\text{Univ Univ} &= \$(\text{T } \$(\text{S } \$\text{Univ}) \text{ Univ}) \end{aligned}$$

The strength annotation on the first predicate argument in the second line above forces the pattern predicate to assert definedness of the subpattern $(\text{S } x)$.

4.3.2 The domain of a pattern

Informally, the *domain* of a pattern is the set of terms that match the pattern. The criterion for matching patterns in Haskell is complicated somewhat by the possibility that a control-disabled subpattern may be embedded into a normally stricter host pattern. In program execution, the match of a control-disabled pattern that is embedded in a case branch is deferred, pending evaluation of the body of the case branch. The match is dynamically performed only if the body is strict in a variable that occurs in the pattern. When a match failure occurs during a deferred pattern match, the match failure is unrecoverable.

We define the domain of a pattern as a predicate characterizing the set of terms matching the pattern in a non-deferred match.

Definition 4.3.2.1: $\text{Dom}(p)$, is the predicate defined by applying the predicate pattern constructor derived from a pattern, p , to a list of *Univ* predicates.

$$\text{Dom}(p) =_{\text{def}} \pi(p) \text{ Univ } \cdots \text{ Univ}$$

The domain predicate of a pattern is calculated by:

$$\text{dom } p = \text{fst } (\text{skeleton } p \text{ (_ } \rightarrow \text{ Univ) True})$$

□

Notice that $\text{Dom}(p)$ is either *Univ* or it is a strong predicate.

The formula $\neg Dom(p)$ asserts that a term fails to match p or is undefined. A strengthened domain predicate disjoined with its strong complement is in effect, a partial definedness predicate. Any term that satisfies either $Dom(p)$ or $\$ \neg Dom(p)$ is well defined in every subterm necessary to evaluate a control-enabled match with the pattern p .

4.3.3 Properties of case branches

A case branch has a pair of properties, one that it exhibits when a case discriminator matches its pattern and another that characterizes its behavior when pattern-matching fails:

$$(3) \quad \frac{\Gamma, x_1 :: P_1, \dots, x_n :: P_n \vdash_{\mathcal{P}} t :: Q}{\Gamma \vdash_{\mathcal{P}} \{p \rightarrow t\} :: \pi(p) P_1 \cdots P_n \rightarrow \$Just\ Q}$$

where $[x_1, \dots, x_n] = fringe\ p$, and

$$(4) \quad \Gamma \vdash_{\mathcal{P}} \{p \rightarrow t\} :: \$ \neg Dom(p) \rightarrow \$Nothing$$

4.3.4 Properties of case expressions

The basic rules for a case expression are those for a single case branch:

$$(5) \quad \frac{\Gamma \vdash_{\mathcal{P}} d :: \pi(p)[P_1, \dots, P_k] \quad \Gamma \vdash_{\mathcal{P}} br :: \pi(p) P_1 \cdots P_n \rightarrow \$Just\ Q}{\Gamma \vdash_{\mathcal{P}} \text{case } d \text{ of } \{br\} :: \$Just\ Q}$$

$$(6) \quad \frac{\Gamma \vdash_{\mathcal{P}} d :: \$ \neg Dom(p)}{\Gamma \vdash_{\mathcal{P}} \text{case } d \text{ of } \{p \rightarrow t\} :: \$Nothing}$$

The following rules account for a **case** expression in which multiple case branches are listed.

$$(7) \quad \frac{\Gamma \vdash_{\mathcal{P}} \text{case } d \text{ of } \{br\} :: \$Nothing \quad \Gamma \vdash_{\mathcal{P}} \text{case } d \text{ of } \{brs\} :: \$Q}{\Gamma \vdash_{\mathcal{P}} \text{case } d \text{ of } \{br, brs\} :: \$Q}$$

where Q has the kind *Maybe_Pred*.

$$(8) \quad \frac{\Gamma \vdash_{\mathcal{P}} \text{case } d \text{ of } \{br\} :: \$Just\ P}{\Gamma \vdash_{\mathcal{P}} \text{case } d \text{ of } \{br, brs\} :: \$Just\ P}$$

Two rules relate a property of a Haskell term of kind *Maybe_Pred* to a property of kind *Pred*.

$$(9) \quad \frac{\Gamma \vdash_{\mathcal{P}} t :: \$ (Just\ P)}{\Gamma \vdash_{\mathcal{P}} t :: P} \quad \frac{\Gamma \vdash_{\mathcal{P}} t :: \$Nothing}{\Gamma \vdash_{\mathcal{P}} t :: Univ}$$

These rules allow a property of a **case** expression to be propagated to a context that expects a property of kind *Pred*. When a pattern match can be proven to fail, the concluded property, $t :: Univ$, provides no specific information.

Figure 6 presents a sample derivation demonstrating how P-logic distinguishes pattern-matching success and failure. The second proof involves a **case** expression which generates a pattern match failure. Interestingly, the strongest property derivable of this expression in P-logic is *Univ*.

$$\begin{array}{c}
\frac{}{\vdash L :: \text{Univ}} \quad \frac{}{\vdash R :: \$R} \quad (1) \quad \frac{}{x :: \text{Univ}, y :: \$R \vdash y :: \$R} \quad (3) \\
\frac{}{\vdash (TLR) :: \$(T \text{Univ } \$R)} \quad (1) \quad \frac{}{\vdash \{(T \sim (S x) y) \rightarrow y\} :: \$(T \text{Univ } \$R) \rightarrow \$\text{Just } \$R} \quad (5) \\
\hline
\frac{}{\vdash \text{case } (TLR) \text{ of } \{(T \sim (S x) y) \rightarrow y\} :: \$\text{Just } \$R} \quad (9) \\
\vdash \text{case } (TLR) \text{ of } \{(T \sim (S x) y) \rightarrow y\} :: \$R \\
\\
\frac{}{\vdash (TLR) :: \$\neg(S \text{Univ})} \quad (2) \\
\frac{}{\vdash \text{case } (TLR) \text{ of } \{(S x) \rightarrow x\} :: \$\text{Nothing}} \quad (6) \\
\hline
\vdash \text{case } (TLR) \text{ of } \{(S x) \rightarrow x\} :: \text{Univ} \quad (9)
\end{array}$$

Fig. 6. Distinguishing Pattern Match Success from Failure in the Logic
(Numbers refer to the rule that applies at each step.)

5 Conclusion

We have presented a two succinct formalisms that specify the reduction semantics of Haskell pattern-matching, a surprisingly complex aspect of the language. The deferred matching that is required of control-disabled (\sim) patterns is of particular interest. In the dynamic semantics, expressed as a meta-language program in Haskell, a deferred pattern match is embedded in a continuation that substitutes the value `bottom` in the bindings of pattern variables in case the match fails. In the verification logic, a pattern predicate is calculated from the pattern and the predicates assumed to be necessary to prove a property of the body of a case branch. This also ensures that a proof of properties will discriminate between failure of a deferred match and failure of a normal match.

In retrospect, one may question whether the generalization of control-disabled patterns, allowing (\sim) annotations to be embedded in a pattern in any context, has been worth the complexity that it adds to understanding the operational semantics of Haskell. Nevertheless, in formulating a programming logic for an existing language one must deal with what is there, and that we have done.

The dual development of a denotational semantics and a verification logic for a programming language offers the opportunity to check soundness of the logical inference rules relative to the model provided by the semantics. When the semantics is also given in an executable framework, as is the semantics of Haskell, soundness checking can be partially automated, but a discussion of that topic is left to another paper.

Acknowledgment The authors wish to thank their colleagues on the Programatica project, particularly Jim Hook, Mark Jones and Sylvain Conchon for their encouragement and for numerous discussions on aspects of logic and Haskell semantics.

References

- [1] Jean-Yves Girard, Yves Lafont, and P. Taylor. *Proofs and types*. Cambridge University Press, 1989.
- [2] Carl A. Gunter. *Semantics of Programming Languages: Programming Techniques*. The MIT Press, Cambridge, Massachusetts, 1992.
- [3] William Harrison, Timothy Sheard, and James Hook. Fine Control of Demand in Haskell. In *6th International Conference on the Mathematics of Program Construction, Dagstuhl, Germany (to appear)*, 2002.
- [4] Simon Peyton Jones and editors John Hughes. Report on the programming language Haskell 98. Technical Report YALEU/DCS/RR-1106, Yale University, CS Dept., February 1999.
- [5] Programatica Home Page. www.cse.ogi.edu/PacSoft/projects/programatica. James Hook, Principal Investigator.
- [6] P. Wadler. The essence of functional programming. *19th POPL*, pages 1–14, January 1992.